


CS 347M (Operating Systems Minor)

Spring 2022

Lecture 1: Introduction to Operating Systems

Mythili Vutukuru
CSE, IIT Bombay

What is an operating system?



- Middleware between user programs and system hardware
 - Not user application software but **system software**
 - Example: Linux, Windows, MacOS
- Manages computer hardware: CPU, main memory, I/O devices (hard disk, network card, mouse, keyboard etc.)
 - User applications do not have to worry about system details
- Operating system has **kernel** + system programs
 - Kernel = the core part of the operating system
 - System programs are useful programs to manage system (e.g., program to list all files in a directory "ls")


History of operating systems

theory
applications
systems ✓

- Started out as a library to provide common functionality to access hardware, invoked via function calls from user program
 - Convenient to use OS instead of each user writing code to manage hardware
 - Centralized management of hardware resources is more efficient
- Later, computers evolved from running a single program to multiple processes concurrently
 - Multiple untrusted users must share same hardware
- So OS evolved to become trusted system software providing isolation between users, and protecting hardware
 - Multiple users are isolated and protected from each other
 - System hardware and software is protected from unauthorized access by users

Course overview

- How operating systems work
 - What is the functionality provided by OS to user programs?
 - How is the functionality of an OS implemented?
- Examples from a simple “teaching” OS called xv6
 - Read xv6 code to fully understand each OS concept
- Programming assignments to reinforce theory
 - Write user programs using OS functionality
 - Adding new functionality to OS within xv6
- Better understanding of computer systems in general
 - Write better programs for any application in future, by knowing how OS works



Course logistics

- Live classes during regular lecture hours on MS Teams
 - Recording available on Teams for those who cannot attend
 - Attendance to online is strongly encouraged for better learning
 - May move to hybrid class model later on (COVID rules permitting)
- Course Moodle Page
 - Summary of weekly content taught, PDF slides from class
 - Programming assignments problem statement and submission link
 - Discussion forum to ask doubts after class
- References and Textbooks:
 - Course material archive (videos, slides, practice problems etc.)
<https://www.cse.iitb.ac.in/~mythili/os/>
 - Highly recommended online textbook: Operating Systems: Three Easy Pieces (OSTEP)
<http://pages.cs.wisc.edu/~remzi/OSTEP/>

Evaluation and Grading

- Four modules in the course: processes, memory, concurrency, I/O
- 3 quizzes, one for each module, 10% each
 - End of Jan, end of Feb (midsem week), end of March
 - Online proctored exams
- 3 programming assignments, one for each module, 10% each
 - Due roughly before the corresponding quizzes
 - Evaluation via demo+viva with TAs
- End-semester exam covering the entire syllabus, 40%
 - Online proctored or offline exam (COVID situation permitting)

endsem

Programming assignments

- Very important part of the course, to understand the concepts better
- Will require non-trivial amount of programming
 - Please take this course only if you have an interest/aptitude for programming
- Individual assignments, to be solved and submitted by each of you
- Discussions with classmates is ok, but no showing/sharing/copying of code is allowed, all code you submit must be written by you only
- ★
 - Plagiarism detection will be run over all submissions
 - Any copying detected will result in FR for all involved students
- Demo of test cases, viva to explain your code and answer questions

Clarifying doubts

- If you have questions or doubts on the course contents:
 - Ask in online class (type in chat box, unmute and speak when asked to)
 - Ask on the Moodle discussion forum. You are encouraged to answer questions from your classmates as well
 - In-person office hours when permitted by COVID guidelines
- If you would like to discuss anything with me in person, please send me email mythili@cse.iitb.ac.in
 - Please avoid using personal email for regular doubts that you can ask in class
 - Please do not send private chat messages on Moodle or Teams
- Please DO NOT hesitate to approach me and talk to me for any question or doubt, no doubt is too stupid to ask ☆

Missing exams and assignment deadlines


- If you have a medical or any other emergency before a quiz or a programming assignment deadline:
 - Please get in touch with me early before the deadline (not after), with suitable proof (e.g., medical certificate)
 - Send email to me with request, cc TA for programming assignments
- Suitable help will be provided as best as we can
 - Assignment deadlines can be extended by 1 or 2 days on a case-by-case basis
 - Makeup quiz will be conducted at the end of the semester, covering all modules, for those who missed any one of the quizzes

Rest of this lecture


- Overview of computer hardware
 - Not the focus of this course, only high level understanding is expected
- Introduction to OS functionality

CPU and memory

CPU ~~memory~~ = I/O

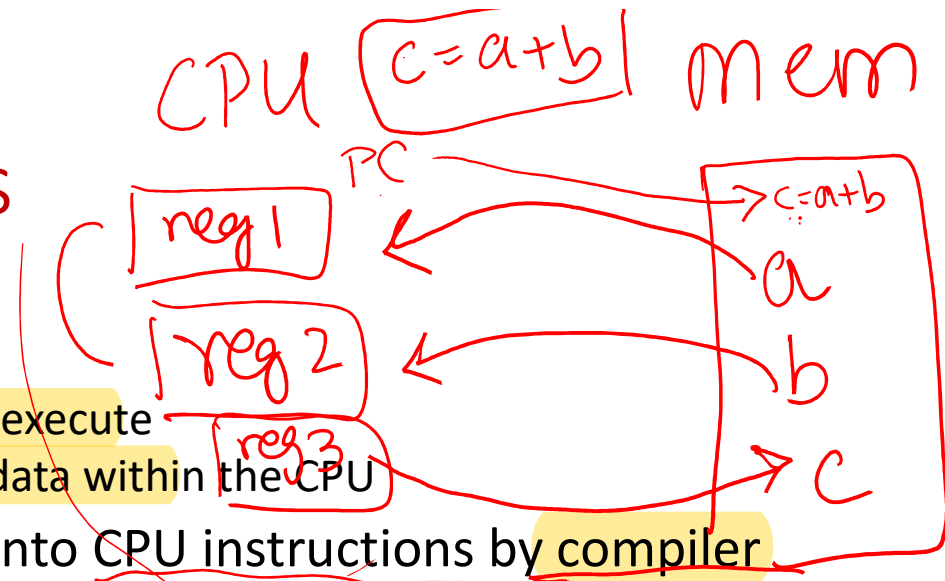


- User **program** = code (instructions for CPU) + data
- Stored program concept
 - User programs stored in main memory or **Random Access Memory (RAM)**
 - Instructions/data occupy multiple contiguous bytes in memory
 - Memory is **byte-addressable**: data accessed via memory address / location / byte#
 - CPU fetches code/data from RAM using memory address, and executes instructions
- CPU runs **processes = running programs**
- Modern CPUs have multiple **CPU cores** for parallel execution
 - Each CPU core runs one process at a time each
 - Modern CPUs have **hyperthreading**, where each core can run more than one process also (OS treats hyper-threading cores also as multiple CPU cores)



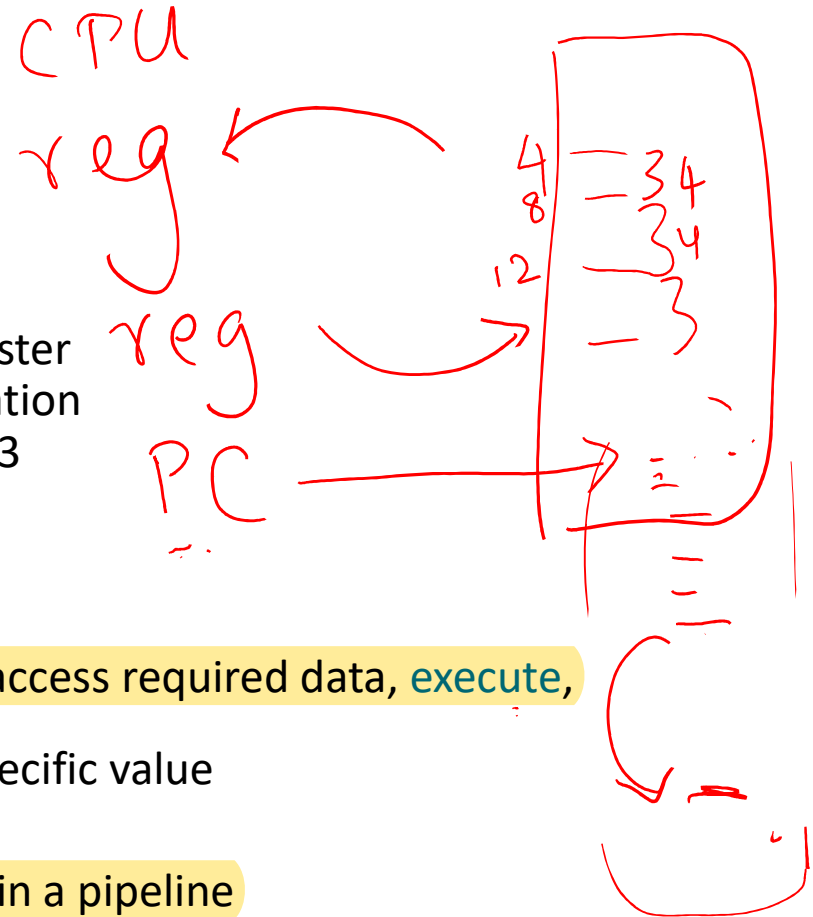
Instructions and registers

- Every CPU has
 - A set of instructions that the hardware can execute
 - A set of registers for temporary storage of data within the CPU
- High level language (C code) translated into CPU instructions by compiler
 - Can directly write assembly language code, but cumbersome
- Instructions and registers defined by ISA = Instruction Set Architecture
 - Specific to CPU manufacturer (e.g., Intel CPUs follow x86 ISA)
- Registers: special registers (specific purpose) or general purpose
 - Program counter (PC) is special register, has memory address of the next instruction to execute on the CPU
 - General purpose registers can be used for anything, e.g., operands in instructions
- Size of registers defined by architecture (32 bit / 64 bit)

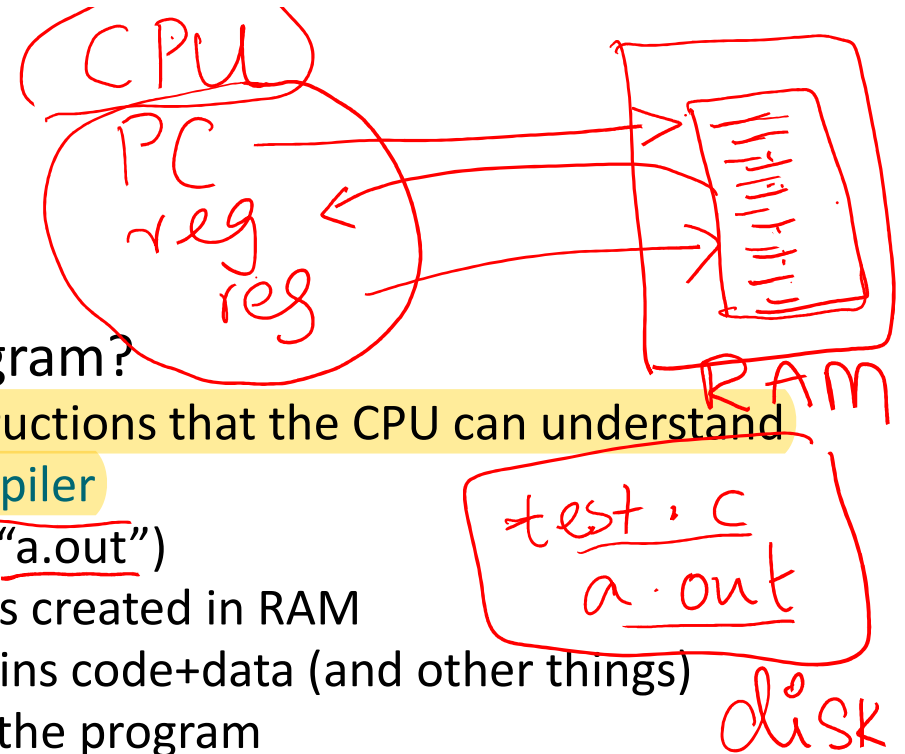


CPU instructions

- Some common examples of CPU instructions
 - Load: copy content from memory location → register
 - Store: copy content from register → memory location
 - Arithmetic operations like add: $\text{reg1} + \text{reg2} \rightarrow \text{reg3}$
 - Logical operations, compare, ...
 - Jump: set PC to specific value
- Simple model of CPU
 - Each clock cycle, **fetch** instruction at PC, **decode**, access required data, **execute**, update PC, repeat
 - PC incremented to next instruction, or jump to specific value
- Many optimizations to this simple model
 - **Pipelining**: run multiple instructions concurrently in a pipeline
 - Many more in modern CPUs to optimize #instructions executed per clock cycle



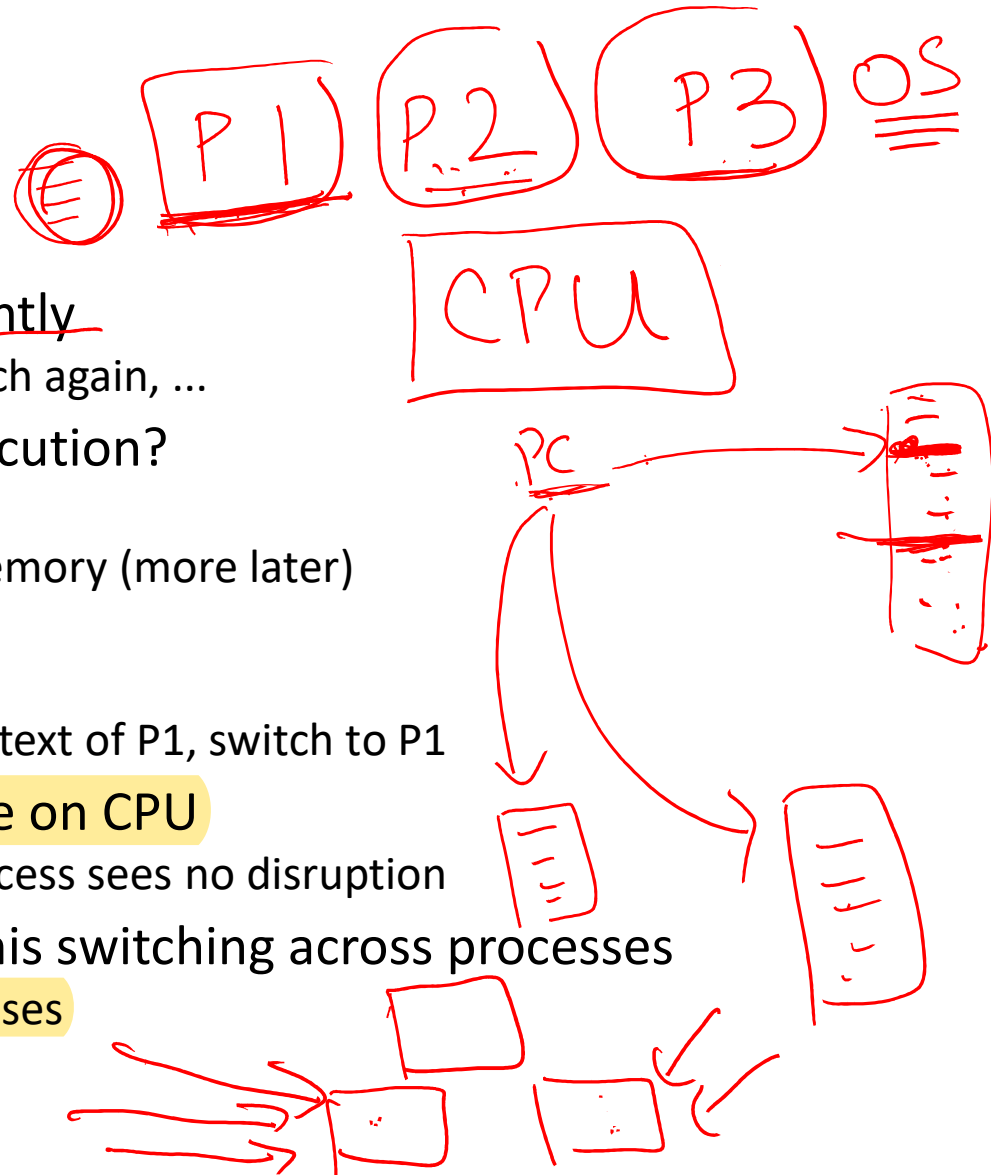
Running a program



- What happens when you run a C program?
 - C code translated into executable = instructions that the CPU can understand
 - Translation done by program called compiler
 - Executable file stored on hard disk (say, "a.out")
 - When executable is run, a new process is created in RAM
 - Memory image of process in RAM contains code+data (and other things)
 - CPU starts executing the instructions of the program
- When CPU core is running a process, CPU registers contain the execution context of the process
 - PC points to instruction in the program, general purpose registers store data in the program, and so on

Concurrent execution

- CPU runs multiple programs concurrently
 - Run one process, switch to another, switch again, ...
- How to ensure correct concurrent execution?
 - Run process P1 for some time
 - Pause P1, save context somewhere in memory (more later)
 - Load context of P2 from memory
 - Run P2 for some time
 - Pause P2, save context of P2, restore context of P1, switch to P1
- Every process thinks it is running alone on CPU
 - Saving and restoring context ensures process sees no disruption
- Operating System (OS) takes care of this switching across processes
 - OS virtualizes CPU across multiple processes




Interrupt handling




- In addition to running user programs, CPU also has to handle external events (e.g., mouse click, keyboard input)
- Interrupt = external signal from I/O device asking for CPU's attention
- How are interrupts handled?
 - CPU is running process P1 and interrupt arrives
 - CPU saves context of P1, runs code to handle interrupt (e.g., read keyboard character)
 - Restore context of P1, resume P1
- Interrupt handling code is part of OS
 - CPU runs interrupt handler of OS and returns back to user code

Isolation



- How to protect processes from one another?
 - Can one process mess up the memory or files of another process?
- Modern CPUs have mechanisms for isolation
- Privileged and unprivileged instructions
 - Privileged instruction = access to sensitive information (e.g., hardware)
 - Regular instructions (e.g., add) are unprivileged
- CPU has multiple modes of operation (Intel x86 CPUs run in 4 rings)
 - Low privilege level (e.g., ring 3) only allows unprivileged instructions
 - High privilege level (e.g., ring 0) allows privileged instructions also
- User code has unprivileged instructions, runs at low privilege level
 - CPU does not execute privileged instructions when in unprivileged user mode
- OS code has privileged instructions, runs at high privilege level
- When user program wants to do privileged operations, it must ask OS
 - CPU shifts to high privilege level, runs OS code, returns to low privilege, back to user code

CPU caches




- CPU must access memory to fetch instructions, load data into registers
 - But main memory (DRAM) is very slow (100s of CPU cycles)
 - CPU cannot do useful work while waiting for memory
- To avoid many memory accesses, CPU stores recently accessed instructions and data in CPU caches
 - Multi-level cache hierarchy, some private to cores, some common
 - Example: private L1, L2, common last level cache (LLC or L3)
 - Can be separate for instructions and data, or common (e.g., L1 is separate)
- Caches have low access latency (tens of CPU cycles), faster than DRAM but smaller in size, more expensive
 - Can only store most recently used instructions and data

RAM



Memory hierarchy



- Hierarchy of storage elements which store instructions and data
 - CPU registers (small number, <1 nanosec)
 - CPU caches (few MB, 1-10 nanosec)
 - Main memory or RAM (few GB, ~100 nanosec)
 - Hard disk (few TB, ~1 millisecc)
- Hard disk is non-volatile storage, rest are volatile
 - Hard disk stores files and other data persistently
- As you go down the hierarchy, memory access technology becomes cheaper, slower, less expensive

