


CS 347M (Operating Systems Minor)

Spring 2022

Lecture 2: Processes

Mythili Vutukuru
CSE, IIT Bombay

Recap: System Hardware




- Users write programs, compile into set of instructions (executable)
- When a **program** is run, a process is created
 - Code and data of program loaded into memory
 - CPU runs instructions of the program one by one
- CPU capable of executing specific set of instructions, using registers for temporary storage of data (**instruction set architecture = ISA**)
- Modern CPUs have multiple CPU cores
 - Each CPU core can run one process at any time, but can concurrently handle multiple processes by timesharing CPU core across processes
 - OS schedules processes on CPU cores, switches between processes on a core



Recap: Memory hierarchy


- Hierarchy of storage elements which store instructions and data
 - CPU registers (small number, <1 nanosec)
 - CPU caches (few MB, 1-10 nanosec)
 - Main memory or RAM (few GB, ~100 nanosec)
 - Hard disk (few TB, ~1 millisec)
- Hard disk is non-volatile storage, rest are volatile
 - Hard disk stores files and other data persistently
- As you go down the hierarchy, memory access technology becomes cheaper, slower, less expensive
- Code data of all active processes stored in RAM
 - Recently used code/data may be available in CPU caches

I/O devices



- A computer system: processor (CPU), main memory, I/O devices connected by system bus
 - Some I/O devices also on special I/O buses (like USB)
 - **Bus**: a set of wires carrying data between components
- Examples of I/O devices
 - Hard disk: stores information in blocks persistently
 - Network Interface Card (NIC): streams information from external machines
 - Keyboard, mouse: input stream from user
 - Display monitor: stream output to user
- I/O takes a long time compared to CPU, so CPU gives command for I/O and device raises interrupt when I/O completes

Device controller and device driver

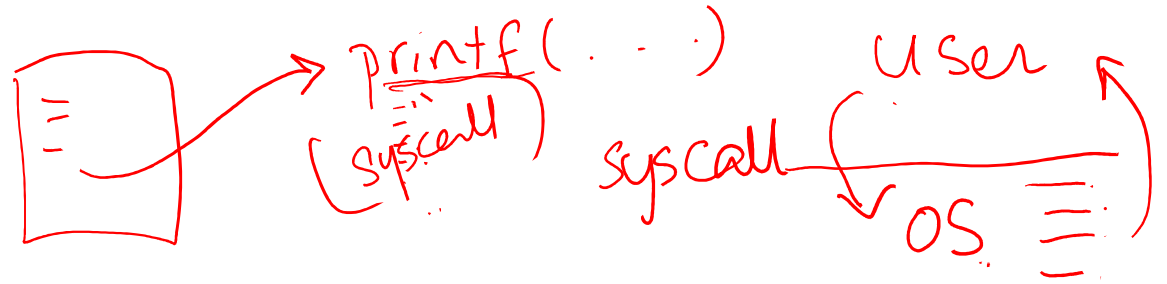


- I/O device is managed by a **device controller**
 - Microcontroller which communicates with CPU/memory over bus
- Device specific knowledge required to correctly communicate with device controller to handle I/O operations
 - Done by special software called device driver
 - Part of operating system code
- Functions performed by device driver
 - Initialize I/O devices
 - Start I/O operations, give commands to device (e.g., read data from hard disk)
 - Handle interrupts from device (e.g., disk raises interrupt when data is ready)

OS manages system hardware for users

- Create and run processes on a system ✓
- Allocate memory to processes ✓
- Schedule processes on CPU cores ✓
- Switch between processes to timeshare CPU ✓
- Give commands to I/O devices to perform I/O operations ✓
- Handle interrupts from I/O devices ✓

System calls

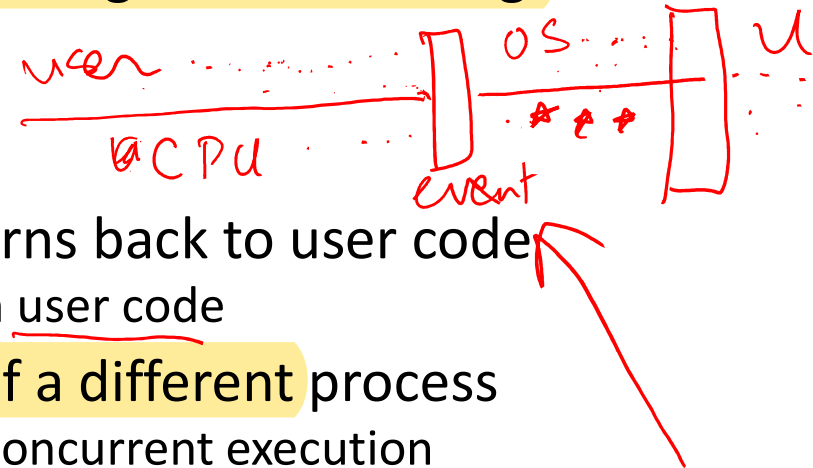


- When user program requires a service from OS, it makes a system call
 - Example: Process makes system call to read data from hard disk
 - Why? User process cannot run privileged instructions that access hardware
 - CPU jumps to OS code that implements system call, and returns back to user code
- System calls supported by an OS form the API to user programs
 - API = application programmer interface
- POSIX API: standard set of system calls defined for portability
 - User program written on one POSIX-compliant OS will run without change on another POSIX-compliant OS
 - However, program may have to be recompiled if architectures are different
- Normally, user program does not call system call directly, but uses language library functions
 - Example: printf is a function in the C library, which in turn invokes the system call to write to screen

User mode and kernel mode


- Modern CPUs operate at multiple privilege levels. *low high*
- User programs run in unprivileged user mode of CPU
- CPU shifts to privileged kernel mode for running OS code during:

- • Interrupts: external events
- • System calls: user request for OS services
- • Program faults: errors that need OS attention



- OS code executes in kernel mode, and returns back to user code
 - CPU switches back to low privilege level to run user code
- Sometimes, can return back to user code of a different process
 - OS can context switch to another process for concurrent execution

Booting your system




- What happens when you boot up a computer system?
- Basic Input Output System (BIOS) starts to run
 - Resides in non-volatile memory, sets up all other hardware
- BIOS locates the boot loader in the boot disk (hard disk, USB, ..)
 - Simple program whose job is to locate and load the OS
- Boot loader loads OS in memory and sets it up for execution
- CPU starts executing OS code
- OS exposes a terminal / shell / other interfaces to user
- User runs programs, starts processes, which start more processes, ...

\$


Questions?

- So far: introduction to system hardware and role of OS
- Next: deep dive into processes

$c = a + b$
increment (c) → f. inc. ? }
→ OS



What defines a process?




- Every process has a unique process identifier (PID)
- Process occupies some memory in RAM (memory image)
 - Code+data from executable
 - Some memory allocated for runtime use
- The execution context of the process (values of CPU registers)
 - PC has address of instruction of process, some registers have process data
 - Process context is in CPU registers when process is running on CPU
 - Context saved in memory when process is paused, restored when run again
- Ongoing communication with I/O devices
 - Information is maintained about files that are open, ongoing network connections, other active connections to I/O devices

Memory allocation in RAM

- When is memory allocated for process code/data in RAM?
- When OS creates process, memory to store compiled executable allocated in RAM *a.out*
 - Executable contains code (instructions) in the program, global/static variables in program
- Should we allocate memory for local variables, arguments of functions in executable?
 - No, since we do not now if/how many times the function will be called at runtime
- Similarly, malloc is for dynamic memory allocation at runtime, not compile time *a.out*

```
int g;  
  
int increment(int a) {  
    int b;  
    b = a+1;  
    return b;  
}  
  
main() {  
    int x, y;  
    x = 1;  
    y = increment(x);  
    print b X  
    int *z = malloc(40);  
    print g  
}
```


Memory image of a process




- Memory image of a process: code+data of process in memory
 - Code: CPU instructions in the program
 - Compile-time data: global/static variables in program executable
 - Stack and heap for dynamic memory allocation at runtime
- Stack: memory allocation during function call
 - Function arguments, local variables etc. pushed onto stack (allocated) at start of function, popped (freed) when function returns
- Heap: dynamic memory allocation by malloc and related function *new*
 - Malloc returns address on heap of allocated memory chunk
- Heap and stack can grow/shrink as process runs, with help of OS
 - Stack pointer CPU register keeps track of top of stack



Example



- Variable "g" allocated at compile time (data)
- Main is also considered function, so variables "x", "y" are allocated on stack when program starts
- When function "increment" is called, variables "a", "b" are allocated on stack
- When malloc is called, 40 bytes allocated on heap
 - Memory address of 40 bytes on heap is stored in pointer variable "z" which is on the stack



```


int g; , h

int increment(int a) {
    int b;
    b = a+1;
    return b;
}

main() {
    int x, y;
    x = 1;
    y = increment(x);
    int *z = malloc(40);
}
    
```

How does OS create a process?

- Allocates memory and creates memory image
 - Loads code, data from disk exe
 - Creates runtime stack, heap
- Opens basic files
 - STD IN, OUT, ERR *initialization*
- Initializes CPU registers
 - PC points to first instruction
- Process starts to run
 - OS steps in as needed for interrupts, system calls, ...



States of a process



- OS manages multiple active processes at the same time. An active process can be in one of the following situations.
- **Running**: currently executing on CPU
 - CPU registers contain context of process
- **Blocked/suspended/sleeping**: process cannot run for some time
 - Example: process has requested data from disk, command issued, but process cannot proceed until the data from disk is available
- **Ready/runnable**: ready to run but waiting for OS scheduler to switch the process in
 - Many processes can be ready but scheduler can only run one on a CPU core
- Context of blocked and ready processes is saved in memory, so that they can continue to run later on

Process State Transitions

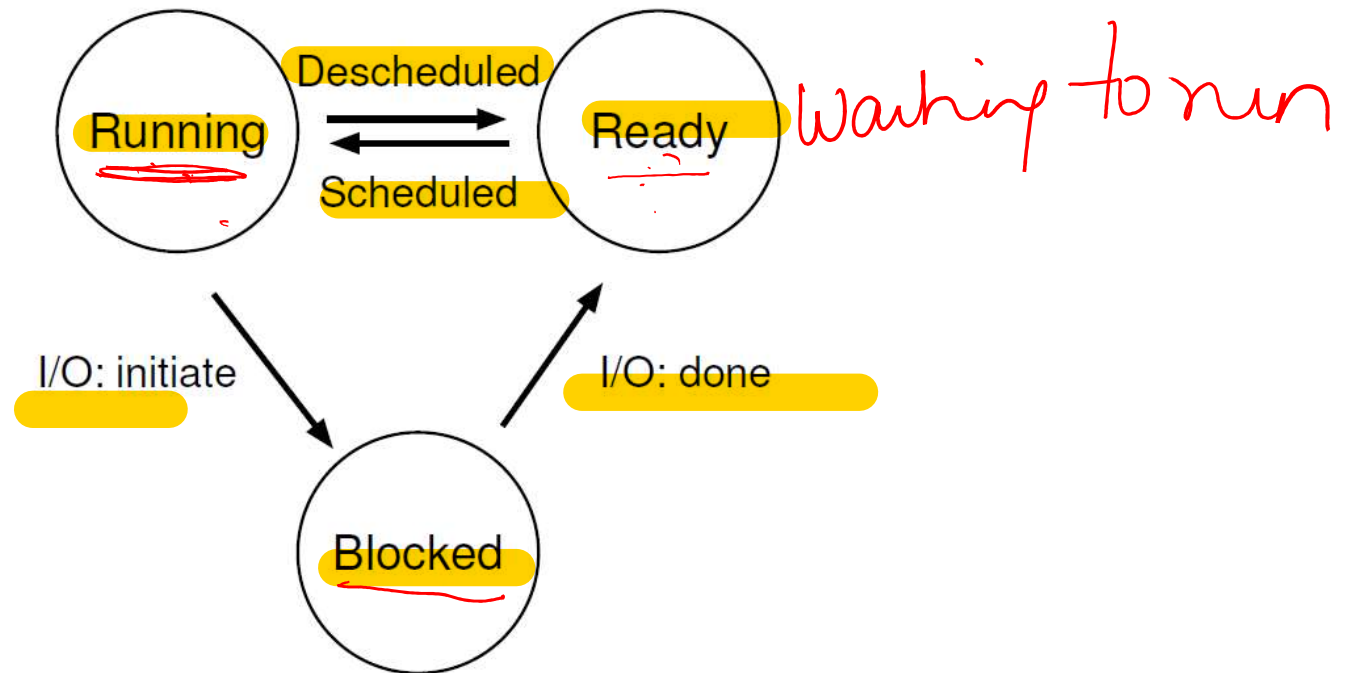
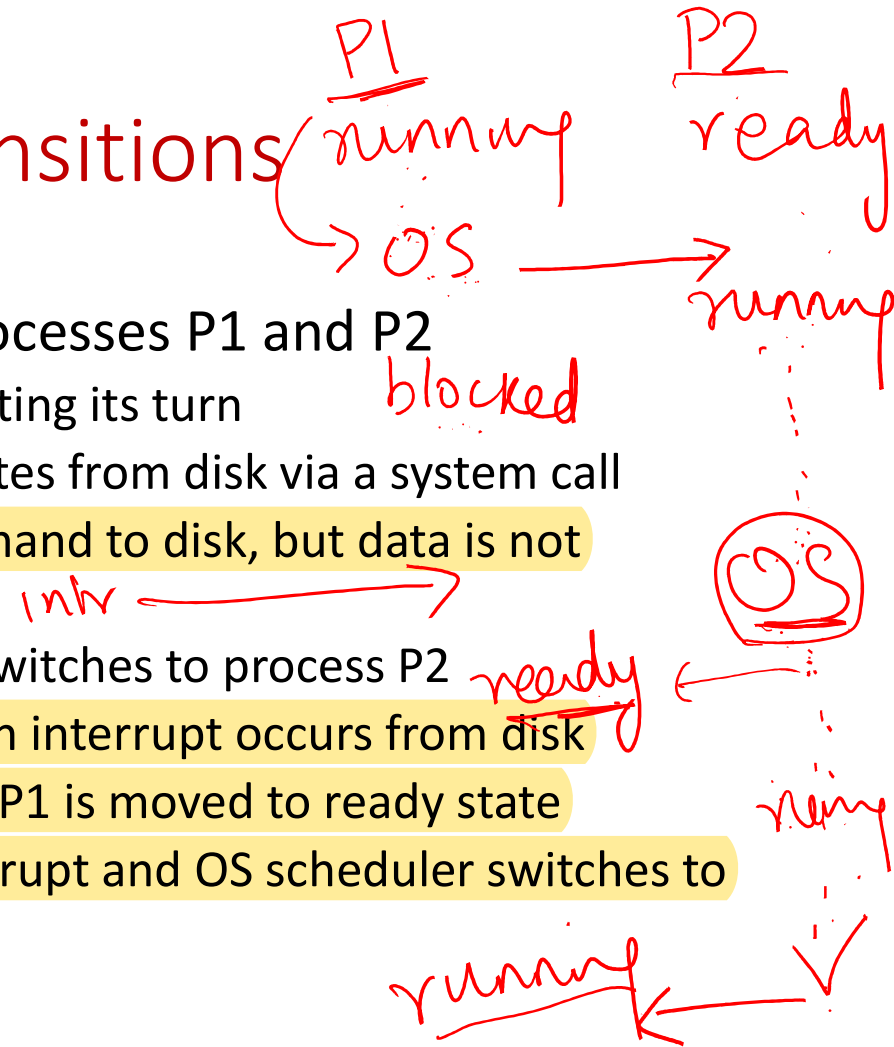


Figure 4.2: Process: State Transitions

Example: process state transitions

- Consider a system that has two user processes P1 and P2
 - Initially P1 is running, P2 is ready and awaiting its turn
 - P1 opens a file and wants to read some bytes from disk via a system call
 - OS handles the system call and gives command to disk, but data is not available immediately
 - Process P1 is moved to blocked state, OS switches to process P2
 - Process P2 runs for some time, and then an interrupt occurs from disk
 - CPU jumps to OS which handles interrupt, P1 is moved to ready state
 - OS can continue to run P2 again after interrupt and OS scheduler switches to ready process P1 later on after some time




Process control block (PCB)

OS data structure


- All information about a process is stored in a data structure called the process control block (PCB)

- Process identifier (PID)
- Process state (running, ready, blocked, terminated, ..)
- Pointers to other related processes (parent, children)
- Saved CPU context of process when it is not running
- Information related to memory locations of a process
- Information related to ongoing I/O communication
- ...

list of all PCB



System call API for processes



- Next lecture: process-related systems calls
 - How to create processes?
 - How to terminate processes?
- fork() creates a new child process
 - All processes are created by forking from a parent
 - The `init` process is ancestor of all processes, which creates shell, which creates other user processes
- exec() makes a process execute a given executable
- exit() terminates a process
- wait() causes a parent to block until child terminates and clean it up

Questions?

- So far: the concept of a process
- Next: processes in xv6

PCB in xv6: struct proc


- PCB is known by different names in different OS
 - struct proc in xv6
 - task_struct in Linux
- Page 23, process structure and process states

```
2334 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
2335
2336 // Per-process state
2337 struct proc {
2338     uint sz; // Size of process memory (bytes)
2339     pde_t* pgdir; // Page table
2340     char *kstack; // Bottom of kernel stack for this process
2341     enum procstate state; // Process state
2342     int pid; // Process ID
2343     struct proc *parent; // Parent process
2344     struct trapframe *tf; // Trap frame for current syscall
2345     struct context *context; // swtch() here to run process
2346     void *chan; // If non-zero, sleeping on chan
2347     int killed; // If non-zero, have been killed
2348     struct file *ofile[NOFILE]; // Open files
2349     struct inode *cwd; // Current directory
2350     char name[16]; // Process name (debugging)
2351 };
2352
```

ready

Process table (ptable) in xv6

```
2409 struct {
2410     struct spinlock lock;
2411     struct proc proc[NPROC];
2412 } ptable;
```



- ptable: Fixed-size array of all processes
 - Real kernels have dynamic-sized data structures
- CPU scheduler in the OS loops over all runnable processes, picks one, and sets it running on the CPU

```
2768 // Loop over process table looking for process to run.
2769 acquire(&ptable.lock);
2770 for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2771     if(p->state != RUNNABLE)
2772         continue;
2773
2774     // Switch to chosen process. It is the process's job
2775     // to release ptable.lock and then reacquire it
2776     // before jumping back to us.
2777     c->proc = p;
2778     switchvm(p);
2779     p->state = RUNNING;
```

Process state transition examples in xv6

- A process that needs to sleep (e.g., for disk I/O) will set its state to SLEEPING and invoke scheduler
- A process that has run for its fair share will set itself to RUNNABLE (from RUNNING) and invoke scheduler
- Scheduler will once again find another RUNNABLE process and set it to RUNNING

run for too long

```
2826 // Give up the CPU for one scheduling round.
2827 void
2828 yield(void)
2829 {
2830     acquire(&ptable.lock);
2831     myproc()->state = RUNNABLE;
2832     sched();
2833     release(&ptable.lock);
```

ready

read

```
2873 void
2874 sleep(void *chan, struct spinlock *lk)
2875 {
2876     struct proc *p = myproc();
2877
2878     if(p == 0)
2879         panic("sleep");
2880
2881     if(lk == 0)
2882         panic("sleep without lk");
2883
2884     // Must acquire ptable.lock in order to
2885     // change p->state and then call sched.
2886     // Once we hold ptable.lock, we can be
2887     // guaranteed that we won't miss any wakeup
2888     // (wakeup runs with ptable.lock locked),
2889     // so it's okay to release lk.
2890     if(lk != &ptable.lock){
2891         acquire(&ptable.lock);
2892         release(lk);
2893     }
2894     // Go to sleep.
2895     p->chan = chan;
2896     p->state = SLEEPING;
2897
2898     sched();
```

for (.)

OS code in C or assembly?

- OS is also like any other program run by CPU, but it is the most important program that manages other programs
 - OS code mostly written in a high-level language like C, compiled into executable, loaded at boot time
- But some parts of OS are written directly in assembly language or CPU instructions that the hardware can understand?
 - Why not write everything in C? Not possible to express certain low level actions performed by OS in high level language
- Basic understanding of x86 assembly code required for understanding xv6 OS code in this course

Reading xv6 for this course

- In this course, only a high level understanding of xv6 is expected
 - For theory as well as labs
- Read C code, comments in code, lecture videos and slides
- No need to read and understand assembly, explanations will be provided for assembly code
- However, it will be useful to understand simple assembly

Reference: x86 registers

- General purpose registers: store data during computations (eax, ebx, ecx, edx, esi, edi)
- Pointers to stack locations: base of stack (ebp) and top of stack (esp)
- Program counter or instruction pointer (eip): next instruction to execute
- Control registers: hold control information or metadata of a process (e.g., cr3 has information related to memory of process)
- Segment registers (cs, ds, es, fs, gs, ss): information about segments (related to memory of process)

Reference: x86 instructions

- Load/store: *mov src, dst*
 - *mov %eax, %ebx* (copy contents of eax to ebx)
 - *mov (%eax), %ebx* (copy contents at the address in eax into ebx)
 - *mov 4(%eax), %ebx* (copy contents stored at offset of 4 bytes from address stored at eax into ebx)
- Push/pop on stack: changes esp
 - *push %eax* (push contents of eax onto stack, update esp)
 - *pop %eax* (pop top of stack onto eax, update esp)
- *jmp* sets eip to specified address
- *call* to invoke a function, *ret* to return from a function
- Variants of above (*movw, pushl*) for different register sizes