

CS 347M (Operating Systems Minor)

Spring 2022


Lecture 8: Paging

Mythili Vutukuru
CSE, IIT Bombay


Recap: Paging

- OS allocates memory in fixed size chunks (“pages”, typically 4KB)
- Avoids external fragmentation (no small “holes”)
- Has internal fragmentation (partially filled pages)
- Virtual address space of process has multiple pages

mem image



RAM



Virtual addresses

Recap: Page table

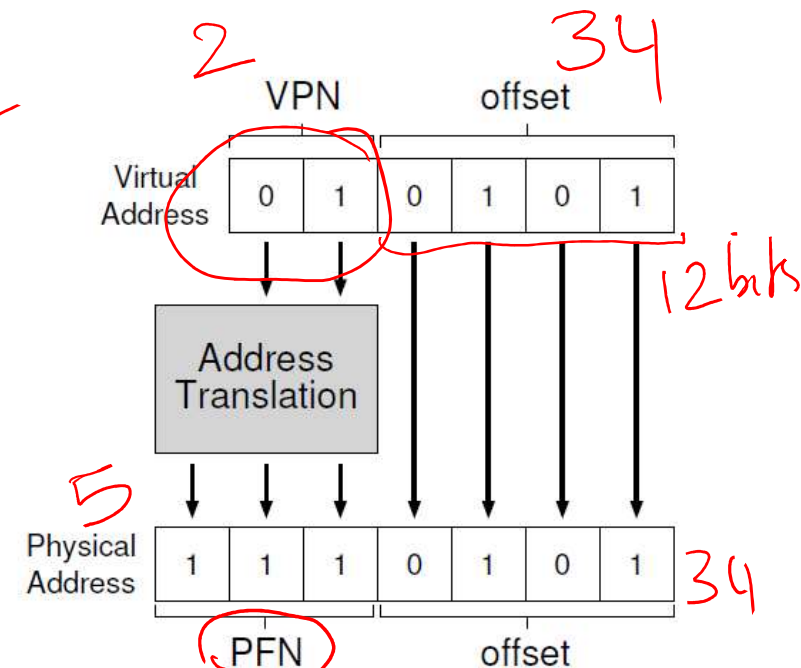


- Per process data structure to help virtual address (VA) to physical address (PA) translation
- Array stores mappings from virtual page number (VPN) to physical frame number (PFN)
 - E.g., VP 0 → PF 3, VP 1 → PF 7
- Part of OS memory (in PCB)
- MMU has access to page table and uses it for address translation
- OS updates page table upon context switch

0	page table: <u>3 7 5 2</u>	page frame 0
16	(unused)	page frame 1
32	page 3 of AS	page frame 2
48	<u>page 0 of AS</u>	page frame 3
64	(unused)	page frame 4
80	page 2 of AS	page frame 5
96	(unused)	page frame 6
112	<u>page 1 of AS</u>	page frame 7
128		

Recap: Address translation in MMU hardware

- Most significant bits of virtual address (VA) give the virtual page number (VPN) $4 \text{ KB} = 2^{12}$
- Least significant bits in VA is offset within page
- Page table array maps VPN to PFN (physical frame number)
- Physical address (PA) is obtained from PFN and offset within a page $2 \rightarrow 5$
- MMU stores (physical) address of start of page table array, not all entries
- MMU accesses suitable entry in page table to translate VA to PA



Handwritten notes and calculations:

- VA: 234 (circled in red)
- page = 100 bytes
- offset: 34
- Physical Address: 534 (calculated as $5 \times 100 + 34$)
- Page # (circled in red)
- Table structure:

0	0 - 99
1	100 - 199
2	200 - 299

4




Recap: What happens on memory access?

- CPU requests code or data at a virtual address, if not found in CPU caches
- **MMU** must translate VA to PA before memory access
 - First, access main memory to read page table entry, translate VA to PA
 - Then, access memory to fetch the actual code/data requested by CPU
- To avoid extra memory access for page table, MMU stores cache of recent VA-PA mappings in **TLB (Translation Lookaside Buffer)**
- If VA found in TLB (TLB hit), memory directly accessed using PA
- If VA not in TLB (TLB miss), MMU accesses memory to read page table, then accesses memory again after computing PA




Structure of page table

page i



- Array of **page table entries**, one per page of process
 - Each page table entry “points to” the physical frame corresponding to the page
- **i-th page table entry (PTE)** contains **physical frame number** and other details (permissions, status, ..) of i-th page of process
 - Valid: is this page in use by process (not all virtual addresses are used by process)
 - Various permission bits (more later)
 - Other status bits: **present, dirty, accessed** (more later)
- Address translation using page table
 - 32 bit virtual address = 20 bit page number + 12 bit offset
 - Use 20 bit page number to index into page table array
 - Find frame number for page number, add offset within page



Storing page table in memory

- What is typical size of page table?

• 32-bit system, $2^{32} = 4\text{GB}$ virtual address space

• Assume page size = $4\text{KB} = 2^{12}$

• Number of pages = $(2^{32}/2^{12}) = 2^{20} = 1\text{M}$

• Assume each page table entry is 4 bytes

• Page table size of one process = 4MB

- How are page tables stored in memory?





• All memory is only allocated in 4KB chunks

- Solution: split page table into pages (much like memory image), use another page table to keep track of original page table!

$$PC = 0 - 2^{32} - 1$$

4GB

$$\frac{4\text{GB}}{4\text{KB}} = 1\text{M}$$



Hierarchical / multi-level page table (1)

2^{20} PTE = 4 bytes

- 4MB page table split into 2^{10} (1K) pages of 4KB each
- Physical frame numbers of 2^{10} pages containing "inner" page tables stored in an outer page table or page directory
 - 4 byte page table entry each, so fits in one page of 4KB
- Page table has two levels
 - **Outer page table (page directory)** has physical frame numbers of 2^{10} "inner" page table pages
 - Each **inner page table** has physical frame numbers of 2^{10} pages of the process virtual address space
- MMU is given the physical address of the outer page directory
 - In case of TLB miss, uses 2-level page table to translate virtual address

Page = 01

VA 0156

4MB

2^{10} 4KB inner page




0-1000
1000-




~~base register~~
ge directory

- Multi-level: MMU stores starting address of outer page directory in page directory base register (CR3 register in x86)



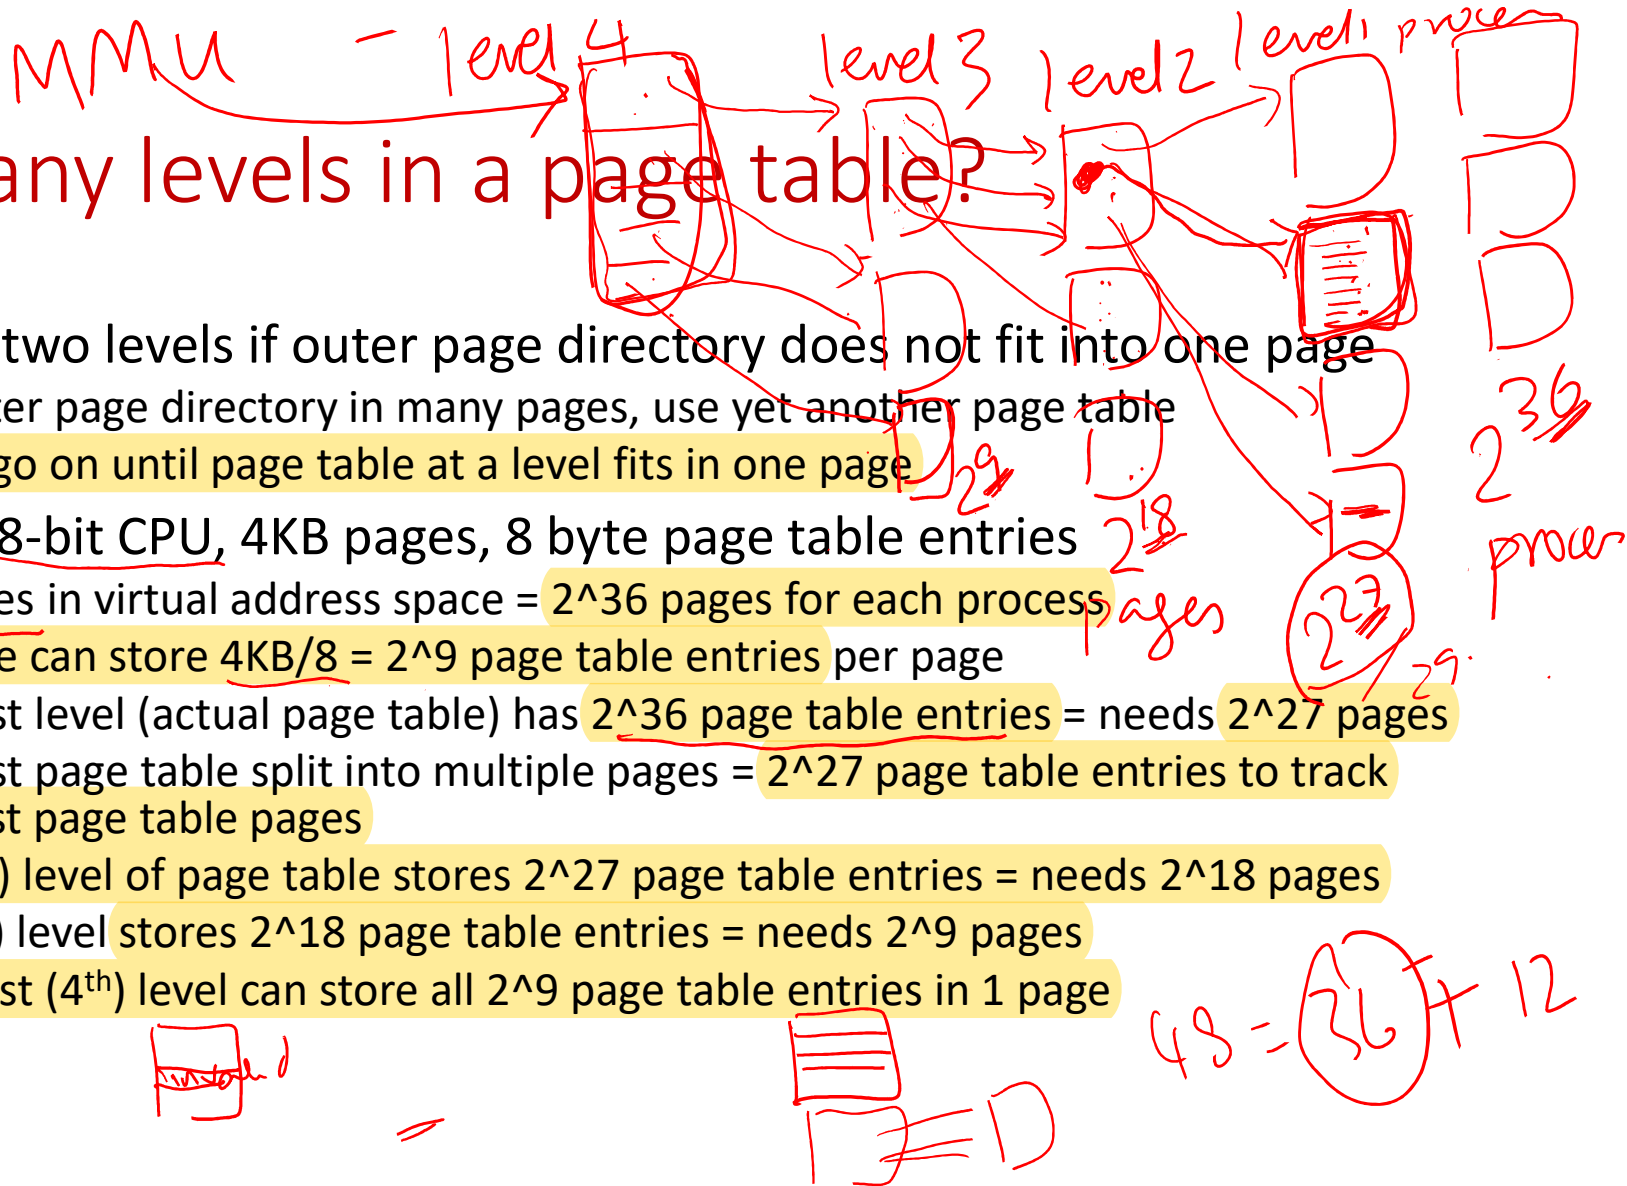
Address translation in 2-level page table

- Virtual address of 32 bits = 20 bit page number + 12 bit offset
 - 20 bits = 10 bit index into page directory, 10 bit index into inner page table
 - Top most 10 bits to index into page directory, identify which one of 2^{10} inner page tables to use, next 10 bits index into inner page table, find one of 2^{10} page table entries, we now have the frame number of a page
 - Computer physical address using frame number and 12-bit offset into page
- MMU "walks" the page table to translate virtual addresses




How many levels in a page table?

- More than two levels if outer page directory does not fit into one page
 - Store outer page directory in many pages, use yet another page table
 - This can go on until page table at a level fits in one page
- Example: 48-bit CPU, 4KB pages, 8 byte page table entries
 - 2^{48} bytes in virtual address space = 2^{36} pages for each process
 - Each page can store $4\text{KB}/8 = 2^9$ page table entries per page
 - Innermost level (actual page table) has 2^{36} page table entries = needs 2^{27} pages
 - Innermost page table split into multiple pages = 2^{27} page table entries to track innermost page table pages
 - Next (2nd) level of page table stores 2^{27} page table entries = needs 2^{18} pages
 - Next (3rd) level stores 2^{18} page table entries = needs 2^9 pages
 - Outermost (4th) level can store all 2^9 page table entries in 1 page



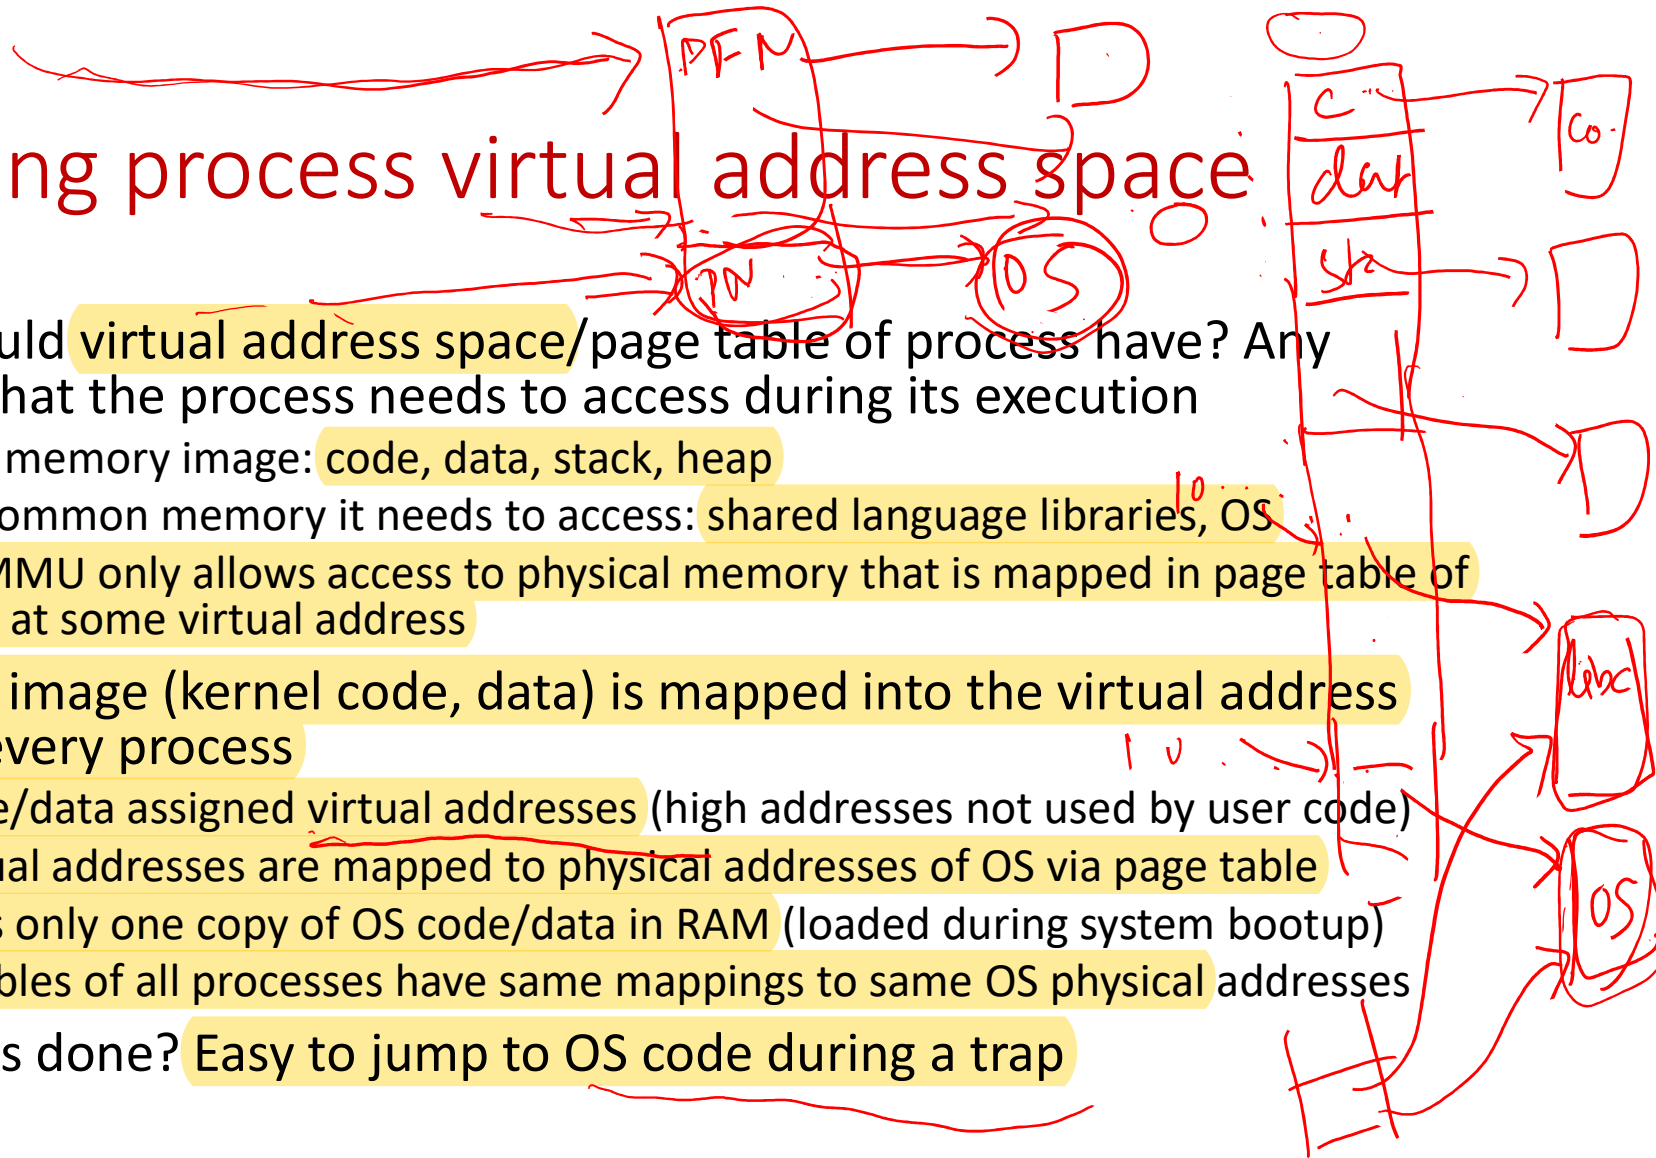
Address translation in multi-level page table

- Example: 48-bit CPU, 4KB pages, 8 byte page table entries
 - 4 level page table required
 - Outmost 4th level page directory has 2^9 entries, containing frame numbers of 3rd level page table pages
 - Each page of 3rd level page table has frame numbers of 2^9 2nd level page table pages
 - And so on....
- How to translate VA to PA?
 - 48-bit VA = 36 bits + 12 bit offset
 - 36 bits = 9 bit offset into each of the 4 levels of page table
- MMU has to access 4 different memory locations for 4 levels of page table, in order to translate one VA to PA
 - MMU page table walks become even longer, TLB hit rate is critical



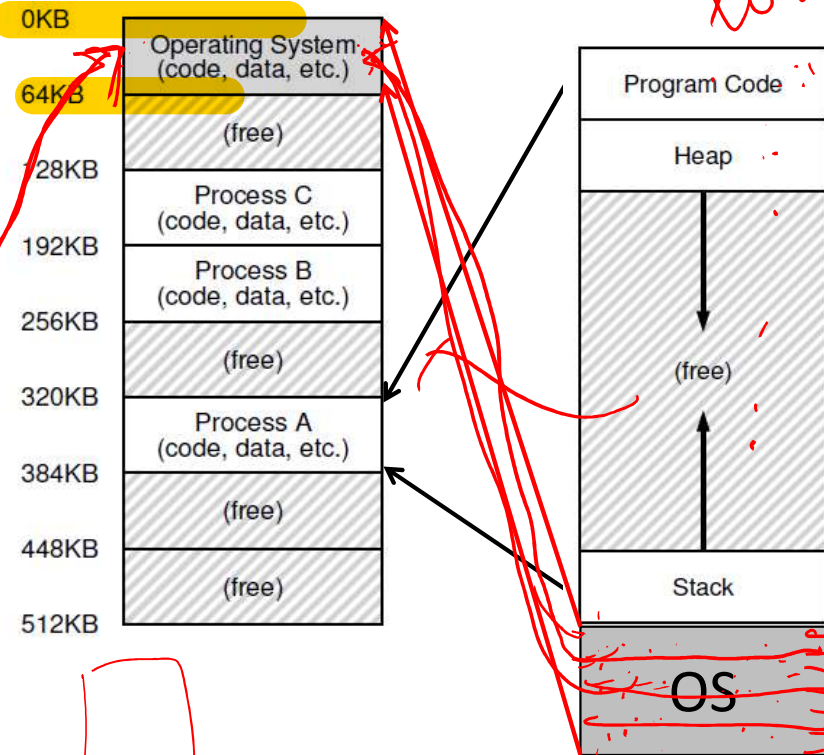
Revisiting process virtual address space

- What should virtual address space/page table of process have? Any memory that the process needs to access during its execution
 - Its own memory image: code, data, stack, heap
 - Other common memory it needs to access: shared language libraries, OS
 - Why? MMU only allows access to physical memory that is mapped in page table of process at some virtual address
- OS binary image (kernel code, data) is mapped into the virtual address space of every process
 - OS code/data assigned virtual addresses (high addresses not used by user code)
 - OS virtual addresses are mapped to physical addresses of OS via page table
 - There is only one copy of OS code/data in RAM (loaded during system bootup)
 - Page tables of all processes have same mappings to same OS physical addresses
- Why is this done? Easy to jump to OS code during a trap

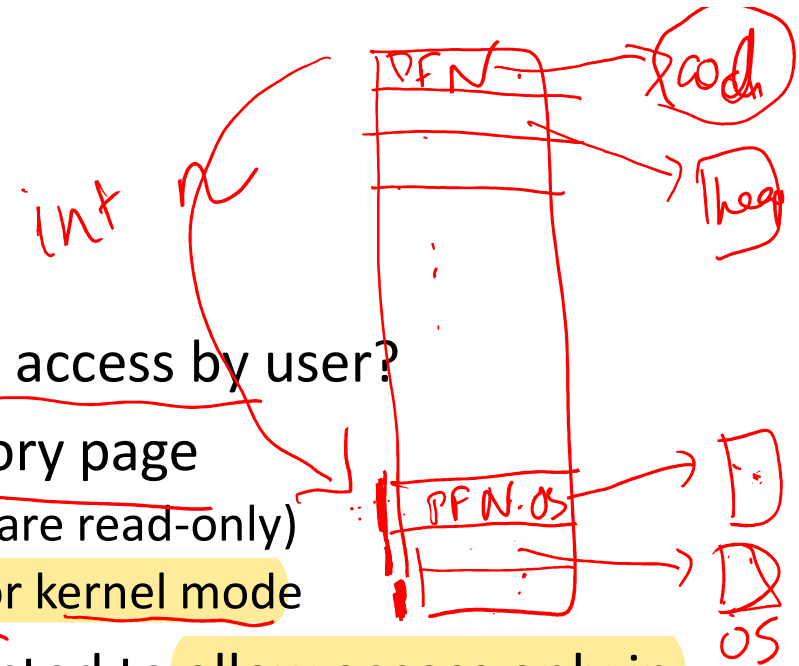


A subtle point: OS is part of address space of every process

- OS is not a separate process with its own address space
- Instead, OS code is part of the address space of every process
- A process sees OS as part of its code (e.g., like a library)
- Page table of every process maps the OS addresses to OS code in main memory
- Only one copy of OS code in memory



Isolation and security



- How is OS code/data protected from illegal access by user?
- Page table has permissions for every memory page
 - Whether read/write or read-only (code pages are read-only)
 - Whether page can be accessed in user mode or kernel mode
- Page table mappings for OS code are protected to allow access only in kernel mode
 - User program in user mode cannot jump to high virtual addresses of OS code
 - CPU in kernel mode (after trap instruction) can access OS code/data
- MMU ensures that user programs cannot access any memory beyond what is visible in the user-accessible part of virtual address space