

CS 347M (Operating Systems Minor)

Spring 2022


# Lecture 9: Demand Paging

Mythili Vutukuru  
CSE, IIT Bombay


# Recap: Virtual addresses and paging



- Instructions and data of a process in memory assigned virtual addresses
  - Starting at 0 for user code (OS code also assigned high virtual addresses)
- Virtual address space of a process divided into fixed size logical pages, stored in a fixed size physical frames in memory
  - Prevents external fragmentation, cannot prevent internal fragmentation
- Page table maps logical page numbers to physical frame numbers
  - One per process, maintained by OS as part of PCB
  - Used by MMU to translate VA to PA when CPU accesses memory



# Demand Paging

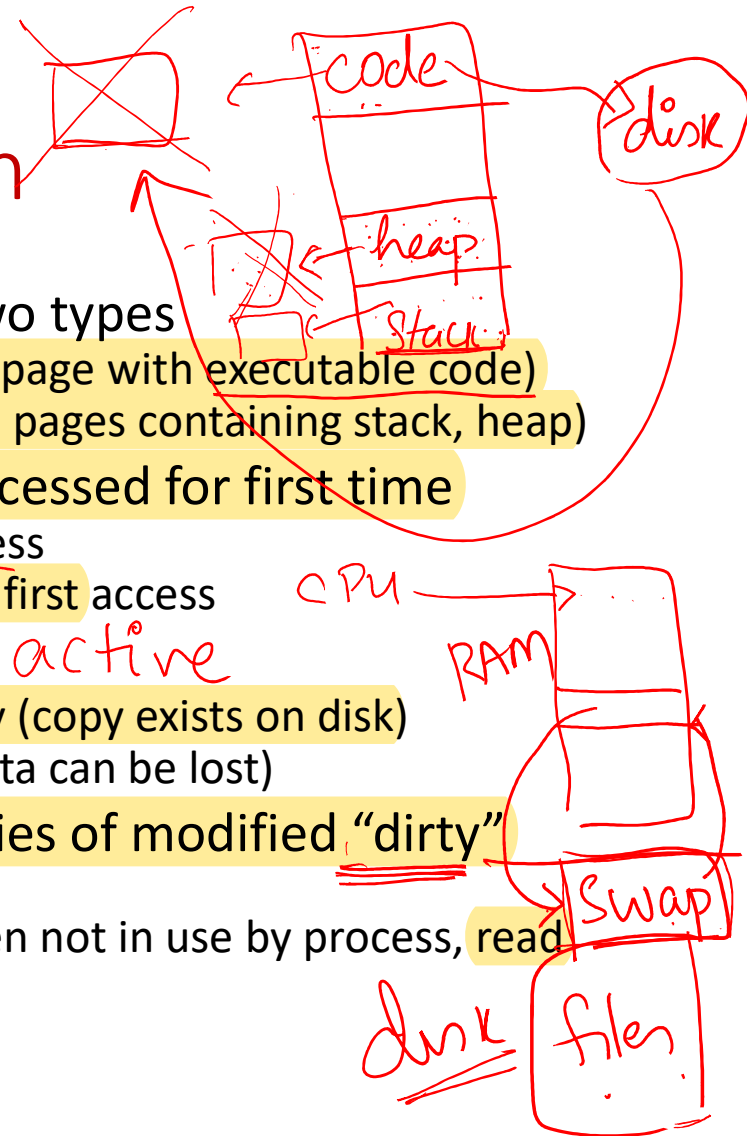


- Are all pages of all active processes always in main memory?
  - Not necessary, as process will not use all of it at once
  - Not possible, with large address spaces
- Modern operating systems provide **virtual memory**
  - Not all logical pages are assigned physical frames
  - OS allocates physical frames to logical pages only on demand, when process accesses the memory contents of the page
  - OS can reclaim some physical memory of process that is not in active use
  - For some pages in page table, physical frame number is not stored, page table entry is marked as "not present"
- Virtual memory of processes can be much more than physical memory in the system, OS overcommits memory


# On-demand memory allocation

- Pages in the memory image of a process are of two types
  - File-backed pages contain data from files on disk (e.g., page with executable code)
  - Anonymous pages are not backed by files on disk (e.g., pages containing stack, heap)
- On-demand memory allocation: allocate when accessed for first time
  - File-backed page content copied from disk on first access
  - Anonymous pages allocated empty physical frames on first access
- How to reclaim unused memory from processes? *active*
  - File-backed page content can be deleted from memory (copy exists on disk)
  - Cannot simply delete content of anonymous pages (data can be lost)
- Swap space: space on hard disk used to store copies of modified "dirty" anonymous pages (different from file storage)
  - Dirty anonymous pages are written to swap space when not in use by process, read back from swap into main memory when required


*how many frames? over commit*



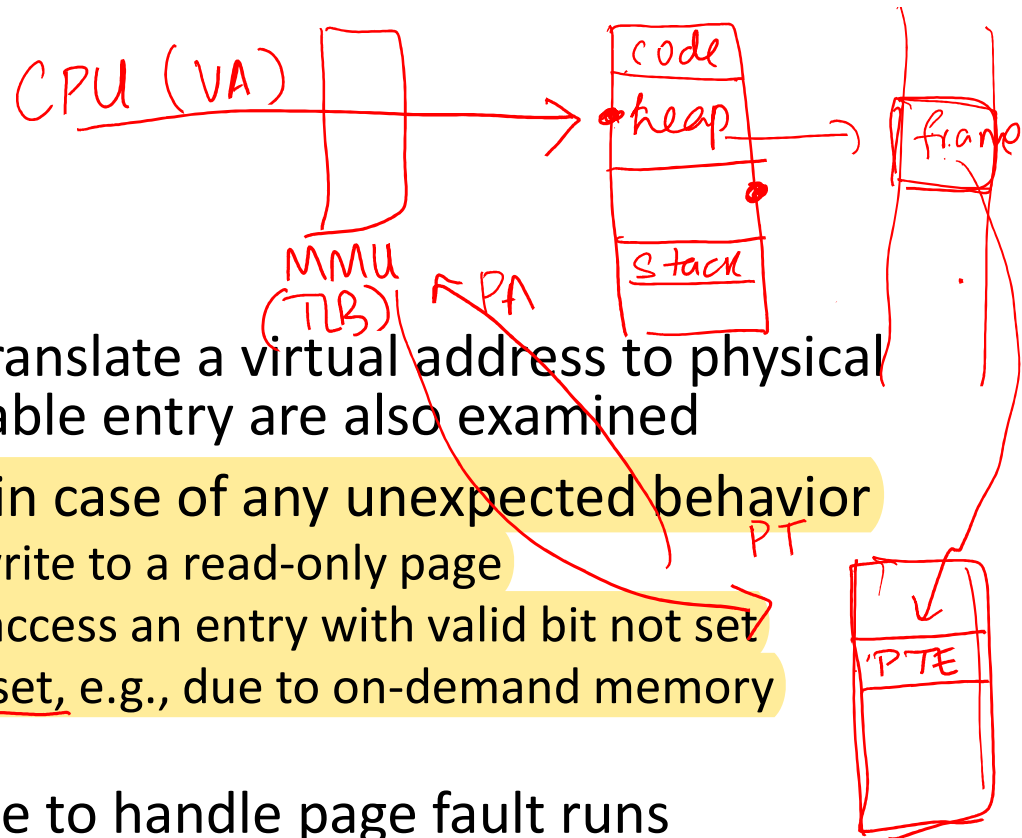
# Information in page table entry



- Page table entry maps page number to physical frame number
- Page table entry also contains several bits to indicate status of page
  - Valid bit indicates if the page is in use by process in its virtual address space (invalid addresses should never be accessed by a process)
  - Present bit indicates if a physical frame number is assigned to the logical page
- Process accesses page with
  - Valid bit not set → illegal memory access (segmentation fault)
  - Valid bit set, present bit not set → OS has not allocated memory yet, or OS has reclaimed memory of this page (content in disk / swap space)
- Other bits set by MMU: dirty bit, accessed bit (more later)




# Page fault handling (1)



- When MMU walks page table to translate a virtual address to physical address, the various bits in page table entry are also examined
- MMU traps to the OS (page fault) in case of any unexpected behavior
  - Illegal access, e.g., process tries to write to a read-only page
  - Invalid access, e.g., process tries to access an entry with valid bit not set
  - Valid bit is set but present bit is not set, e.g., due to on-demand memory allocation not done yet by OS
- Trap instruction is invoked, OS code to handle page fault runs
  - For illegal/invalid accesses, OS may terminate the process when servicing the page fault
  - If the virtual address is valid but not present in physical memory, OS will service page fault by assigning a free frame to the page

traps  
syscalls  
into  
page fault

## Page fault handling (2)




- If page not assigned a frame, OS finds free physical frame to store page
  - OS maintains list of free physical frames to assign during page faults
- Next, OS fills the physical frame with page contents
  - If page is in swap space or file-backed, contents are read from disk into free page
  - Reading the page contents from disk may block the process
- Once physical frame is ready with page content, OS updates page table entry (add new physical frame number), updates MMU, process restarts
  - Hopefully runs correctly this time and MMU doesn't raise trap




TLB flush

## Page fault handling (3)

free list → I → D → D

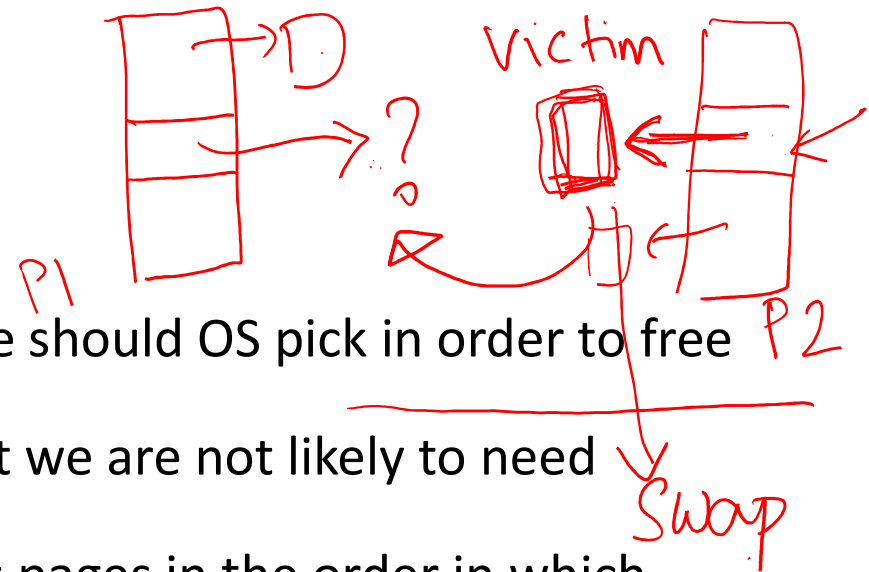


- What if all physical frames are occupied and OS cannot find free physical frame to store page?
  - If there are no free physical frames, OS can evict a victim page (i.e., take away a physical frame assigned to another page) to free up a physical frame
  - Page replacement policy helps OS identify victim pages
  - OS writes victim page contents to swap space (in case of modified “dirty” anonymous page)
  - Page table of victim process updated to delete old mapping
  - Victim page can belong to same process or another process
- Servicing page fault may potentially involve two disk accesses
  - Store old contents of victim page to disk (if dirty)
  - Read new contents of page from disk (if not empty page)





# Page replacement policies



- Page replacement policy: which victim page should OS pick in order to free up a physical frame?
- Goal: Minimize page faults, evict pages that we are not likely to need immediately
- Simple policy: First In First Out (FIFO) evicts pages in the order in which they have been assigned frames
  - May be suboptimal, e.g., the first assigned pages may be important pages that are in use very often, leading to another page fault in near future
- Most commonly used policy: evict the Least Recently Used (LRU) page
  - Page has not been used for sometime now, so less likelihood that it will be immediately used in future
- Ideal optimal policy: replace page not needed for longest time in future (not practical!)



# Example: Optimal policy

- Example: Process accessed 4 pages (0,1,2,3), only 3 physical frames in memory
- First few accesses are cold (compulsory) misses

*references to memory*

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss	3	0, 1, 2
1	Hit		0, 1, 2

*page fault*

Figure 22.1: Tracing The Optimal Policy

# Example: FIFO

- Usually worse than optimal
- Belady's anomaly: performance may get worse when memory size increases!

*reference*

Access	Hit/Miss?	Evict	<del>memory</del> Resulting Cache State
0	Miss		First-in → 0
1	Miss		First-in → 0, 1
2	Miss		First-in → 0, 1, 2
0	Hit		First-in → 0, 1, 2
1	Hit		First-in → 0, 1, 2
3	Miss	0	First-in → 1, 2, 3
0	Miss	1	First-in → 2, 3, 0
3	Hit		First-in → 2, 3, 0
1	Miss	2	First-in → 3, 0, 1
2	Miss	3	First-in → 0, 1, 2
1	Hit		First-in → 0, 1, 2

*memory*

# Example: LRU

- Equivalent to optimal in this simple example
- Works well due to locality of references (recently used pages accessed again with high probability)

*access pattern*

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		LRU→ 0
1	Miss		LRU→ 0, 1
2	Miss		LRU→ 0, 1, 2
0	Hit		LRU→ 1, 2, 0
1	Hit		LRU→ 2, 0, 1
3	Miss	2	LRU→ 0, 1, 3
0	Hit		LRU→ 1, 3, 0
3	Hit		LRU→ 1, 0, 3
1	Hit		LRU→ 0, 3, 1
2	Miss	0	LRU→ 3, 1, 2
1	Hit		LRU→ 3, 2, 1


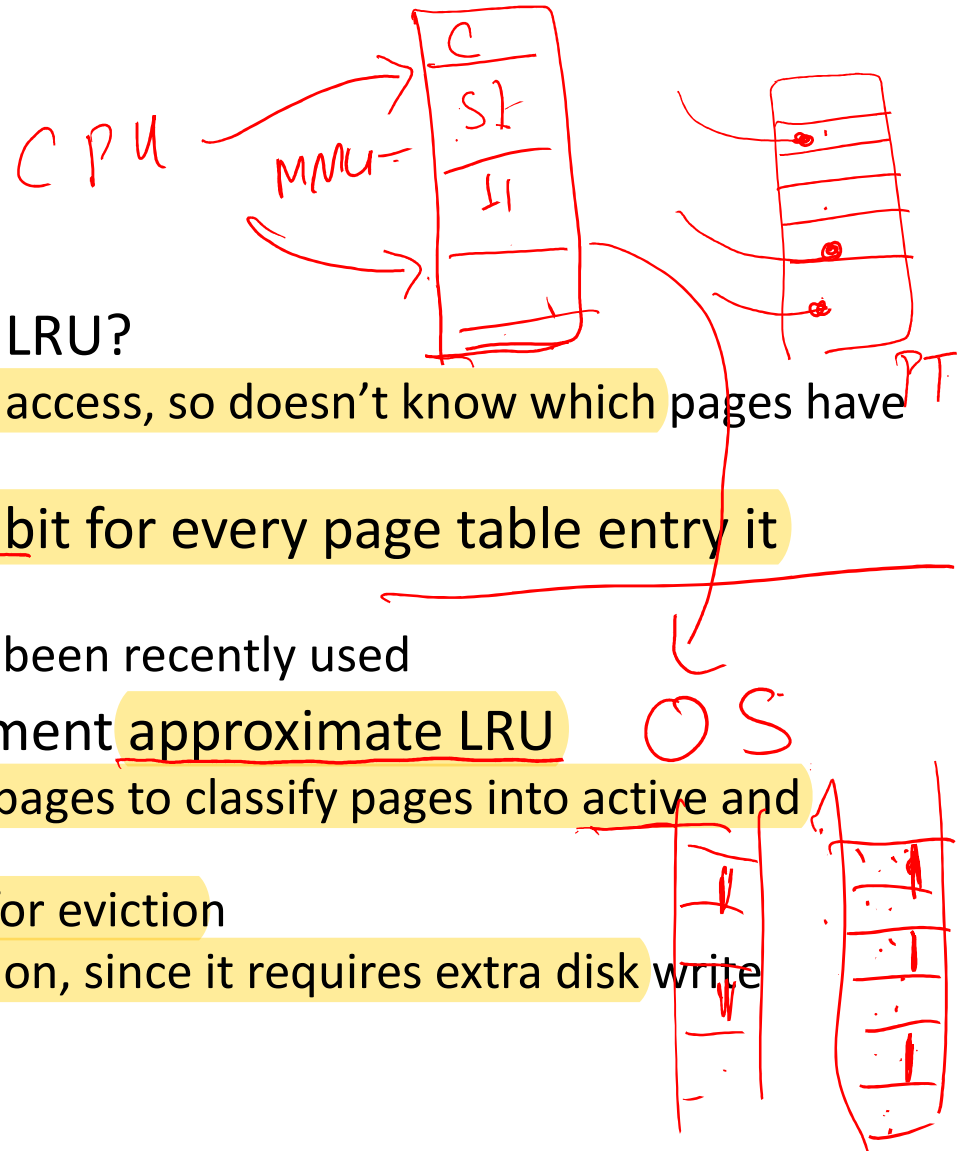



Figure 22.5: Tracing The LRU Policy

# LRU implementation




- How does OS know which page is LRU?
  - OS is not involved in every memory access, so doesn't know which pages have been recently used
- Solution: MMU sets the accessed bit for every page table entry it accesses
  - Accessed bit is set implies page has been recently used
- Modern operating systems implement approximate LRU
  - Periodically, look at accessed bit of pages to classify pages into active and inactive pages
  - Pick pages that have been inactive for eviction
  - May also avoid dirty pages for eviction, since it requires extra disk write



## What happens on a memory access?

- CPU has requested data (or instruction) at a certain memory address
  - CPU checks caches. If hit in **CPU cache**, data is directly available (within ns) RAM
  - If CPU cache miss, CPU has to access main memory (hundreds of ns) 10<sup>-9</sup>
  - MMU checks **TLB**. If TLB hit, the physical address is available, the data is fetched from memory. Otherwise, MMU has to **walk page table**
  - If address is valid and present in page table, permission checks pass, translate address and access memory (MMU adds translation to TLB)
  - For any error in address translation, MMU traps to OS for **page fault**
  - OS may need to read/write from disk to service page fault (few millisec)
- Causes for application slowdown: **CPU cache misses, TLB misses, page faults, ...**

# Thrashing



- How much physical memory should OS assign a process?
- Every process has a working set: frequently used pages in memory image
  - Working set can change from time to time, based on code being executed
  - Working set is usually smaller than total virtual memory of process
  - If memory assigned to process is less than working set, frequent page faults, frequent interruptions to process execution
- Thrashing = system spends too much time servicing page faults and swapping back and forth from disk, and too little time doing useful application work
  - Significant slowdown in application performance will be noticed by users
- Solution: users can reduce working set of processes, OS can terminate some processes or clean up unnecessary memory, ...

