


CS 347M (Operating Systems Minor)

Spring 2022

Lecture 10: Memory allocation in user programs

Mythili Vutukuru
CSE, IIT Bombay



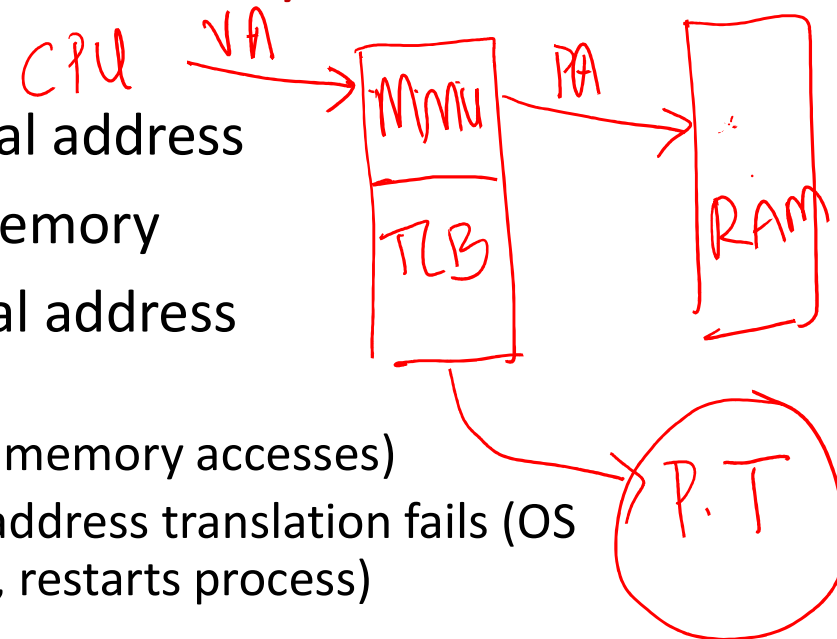
Recap: Virtual memory management



- Code+data of a process assigned virtual addresses
- Fixed size chunks of virtual addresses (pages) mapped to chunks of physical addressss (frames) in main memory
- Page table of process maps page# to frame#
 - Used by MMU for address translation
 - Page table mappings cached in TLB
- Not all valid pages (virtual addresses) assigned frames (physical addresses) by OS always
 - Some pages may be mapped to frames only when used (demand paging)
 - MMU traps to OS when it cannot translate address (page fault)

Recap: What happens on a memory access


- CPU accesses code/data at a certain virtual address
- Check CPU cache, else fetch from main memory
- MMU translates virtual address to physical address
 - If TLB hit, physical address directly available
 - If TLB miss, MMU walks page table (multiple memory accesses)
 - During page table walk, MMU traps to OS if address translation fails (OS handles page fault, updates page table entry, restarts process)
- Main memory accessed using physical address



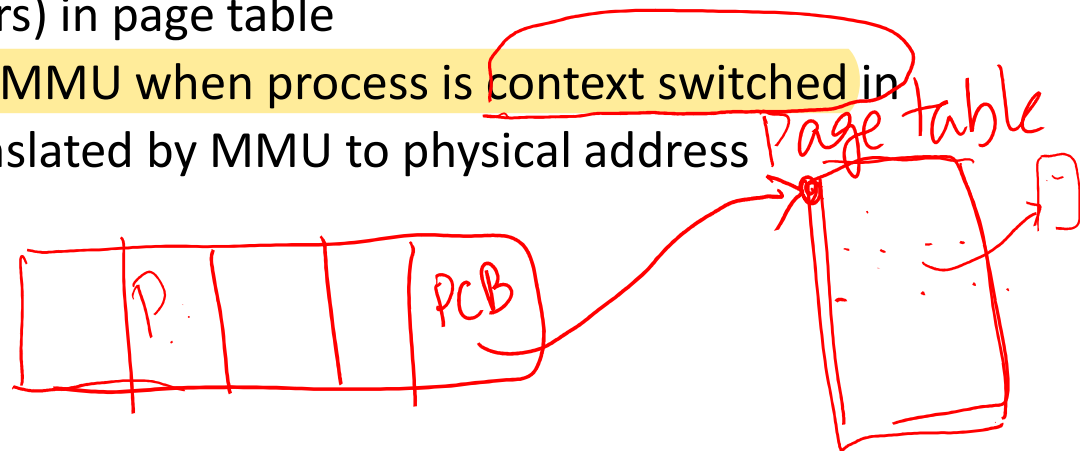
Running a program

P → C
P.T created


- User program is compiled into executable = program code + compile-time data (global/static variables)
- When program is run, OS creates process
 - Allocates some physical frames to store executable code/data, stack, heap, ..
 - Allocates page table, adds mapping from virtual addresses (page numbers) to physical addresses (frame numbers) in page table
 - Pointer to page table provided to MMU when process is context switched in
 - CPU accesses virtual address, translated by MMU to physical address




OS memory



Memory allocation to user programs



- Not all virtual addresses (0 to MAX possible) have corresponding physical addresses
- Not all virtual addresses available are used by processes
- Not all used virtual addresses are assigned physical addresses by OS
- How can process ask for physical memory to be allocated for its valid virtual addresses?
 - Access virtual address which has no memory allocated
 - OS handles page fault and assigns physical memory
 - Keep using memory actively, so that OS doesn't reclaim it
- How can process expand its virtual address space? Ask OS via syscalls



System calls for virtual memory allocation

- Default virtual address space: code+static data, stack, heap, ...
 - Most programs are happy with this, no need to use syscalls to expand
 - Heap expansion taken care of by programming language libraries
- System calls to expand virtual address space
 - The syscalls brk/sbrk are used to allocate page-sized chunks at the end of the data segment of executable (called program break)
 - mmap syscall used to obtain one or more page sized chunks at any unallocated range of virtual addresses
- Syscalls sbrk/brk/mmap used by heap managers to expand heap, can also be used by user programs (instead of malloc) to obtain larger chunks of memory



program break malloc

malloc (s)

malloc

heap



mmap

? ¹⁰ heap

Memory mapping


4KB OS

(int *)

int (0) ...
int (1) ...

```
char *buf = mmap(address, size, ...)  
buf[0] = ...
```

- mmap system call for mapping new memory into address space
 - Used to expand virtual address space
 - Takes starting virtual address, size of memory required as arguments (must be multiple of page size)
 - OS allocates unused virtual addresses, marks page table entries as valid
 - Syscall returns the starting virtual address of the allocation
 - Physical memory corresponding to the virtual addresses allocated on demand
- How is memory-mapped page used? *malloc(5)* *mmap*
 - Allocated memory can be split into smaller chunks by heap manager
 - Allocated memory can be directly used in user programs to store data



char * ptr = malloc (10)
free (ptr)

test.c
malloc



Heap management

syscall
OS



- Heap: one or more pages of memory used for dynamic memory allocation via malloc, managed by programming language libraries
- Heap manager: code within libraries to manage heap
 - Exposes API to allocate/deallocate memory in variable-sized chunks at run time (e.g., malloc, free)
 - Gets page-sized chunks from OS using sbrk/brk/mmap
 - Splits pages into smaller chunks and gives out during malloc
 - Malloc returns the starting virtual address of the free chunk of requested size
 - Freed up memory in heap is reused for future allocations
 - Some language libraries automatically clean unused chunks, some do not (malloc-ed memory must be explicitly freed up by user, else memory leak)



int a
&a



C / C++

malloc
free

Types of heap managers

malloc {
}
implementation


- General-purpose heap managers support variable sized memory allocation using malloc
 - Complex data structures to keep track of variable sized free chunks
 - Frequent calls to malloc/free can slow down user programs
- Some heap managers optimized for fixed size allocation: slab allocators
 - Useful for user applications that allocate memory in fixed sizes
 - Heap memory is divided into fixed size chunks for allocation
 - More efficient than general-purpose variable sized allocation
- Usually managed by programming language libraries, but can be configured/changed by users




free (x)

malloc

Variable sized memory allocation



- How do we design a general purpose heap manager?
- Given a large block of memory (one or more pages), how do we use it to satisfy variable-sized memory allocation requests?
- Why is this hard?
 - Need to track information about variable sized chunks: what sized chunk is allocated where in memory
 - External fragmentation: after memory is allocated and freed-up, many gaps of different sizes throughout heap
 - Need to keep track of various free chunks to reuse for later allocations
 - What if chunk of desired size not found? May need to merge or split chunks



Variable sized allocation: headers

- Consider a simple implementation of `malloc`
- Every allocated chunk has a header with info like size of chunk
 - Why store size? We should know how much to free when `free()` is called
- Header also has some random number (“magic”)
 - Detects if previous chunk has overflowed into this chunk




Figure 17.1: An Allocated Region Plus Header

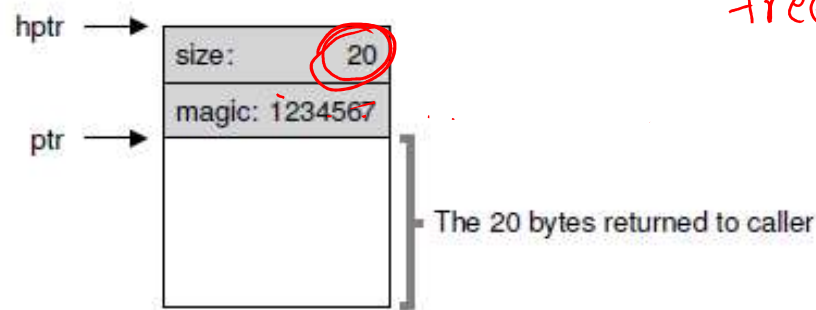



Figure 17.2: Specific Contents Of The Header



free(1004)

Variable sized allocation: free list

- How to store information about which chunks are free?
- Free space can be managed as a linked list of free chunks
 - Pointer to the next free chunk is embedded within the free chunk
- Need to only remember head of list within heap manager
 - Allocations happen from the head




Figure 17.3: A Heap With One Free Chunk

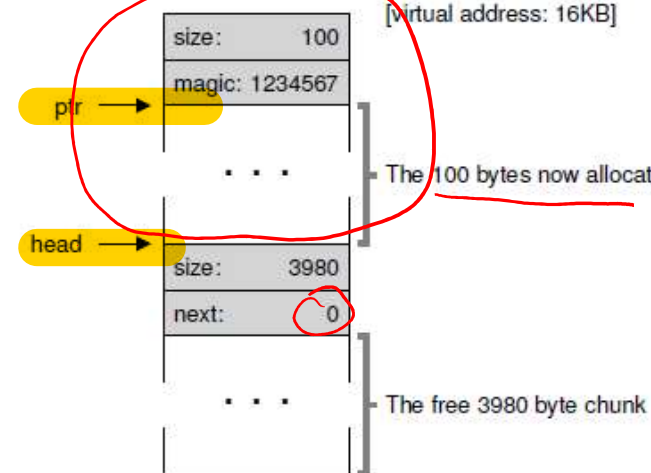
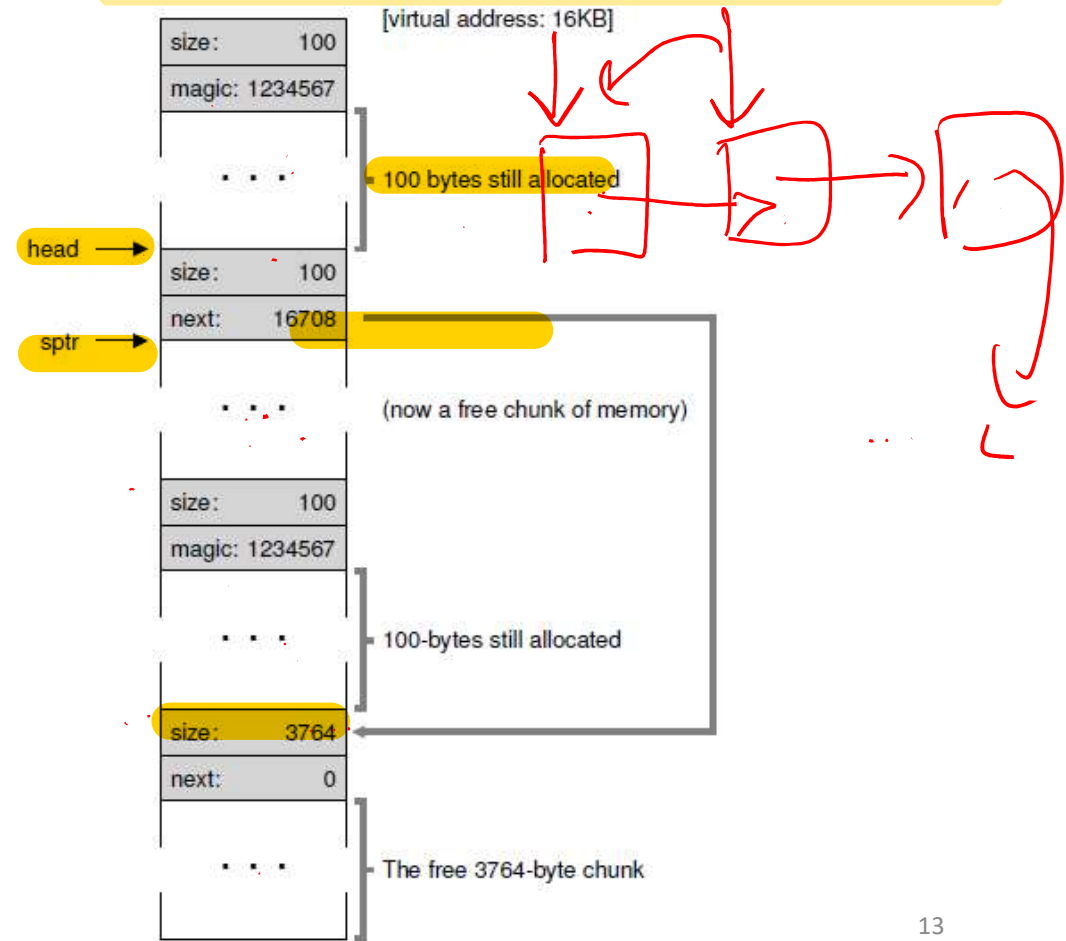


Figure 17.4: A Heap: After One Allocation

Variable sized allocation: external fragmentation

- Suppose 3 allocations of size 100 bytes each happen. Then, the middle chunk pointed to by `sptr` is freed
- What is the free list?
 - It now has two non-contiguous elements
- Free space may be scattered around due to fragmentation
 - Cannot satisfy a request for 3800 bytes even though we have the free space



Variable sized allocation: splitting and coalescing

- Suppose all the three chunks are freed
- The list now has a bunch of free chunks that are adjacent
- A smart algorithm would merge them all into a bigger free chunk
- Must split and coalesce free chunks to satisfy variable sized requests

Linked list

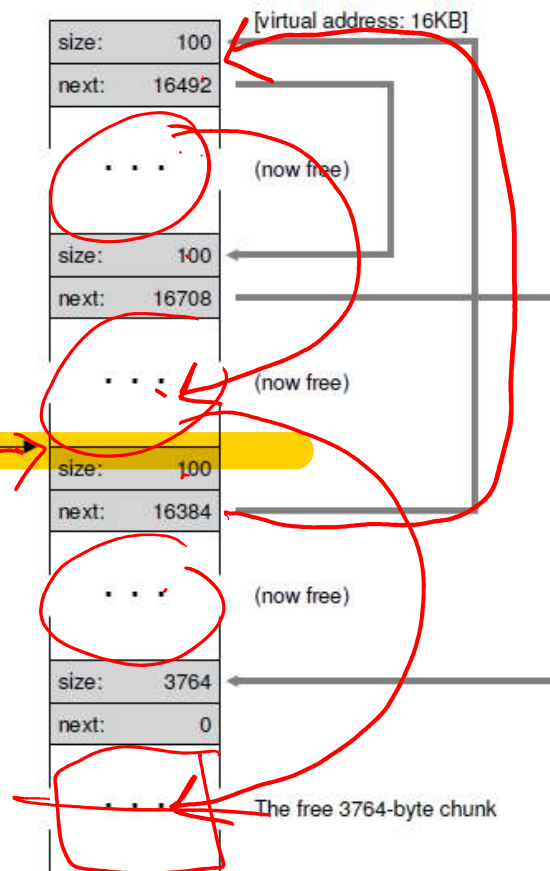



Figure 17.7: A Non-Coalesced Free List

Which free chunk to pick in malloc?


- First fit: allocate first free chunk that is sufficient
- Best fit: allocate free chunk that is closest in size
- Worst fit: allocate free chunk that is farthest in size
- Example, consider this free list, and malloc(15)



- Best fit would allocate the 20-byte chunk



- Worst fit would allocate 30-byte chunk: remaining chunk is bigger and more usable



Buddy allocation for easy coalescing

Slab
buddy

- Allocate memory in size of power of 2
 - E.g., for a request of 7000 bytes, allocate 8 KB chunk
- Why? 2 adjacent power-of-2 chunks can be merged to form a bigger power-of-2 chunk
 - E.g., if 8KB block and its “buddy” are free, they can form a 16KB chunk

