


CS 347M (Operating Systems Minor)

Spring 2022

Lecture 11: Memory management in xv6

Mythili Vutukuru
CSE, IIT Bombay


Virtual memory in xv6



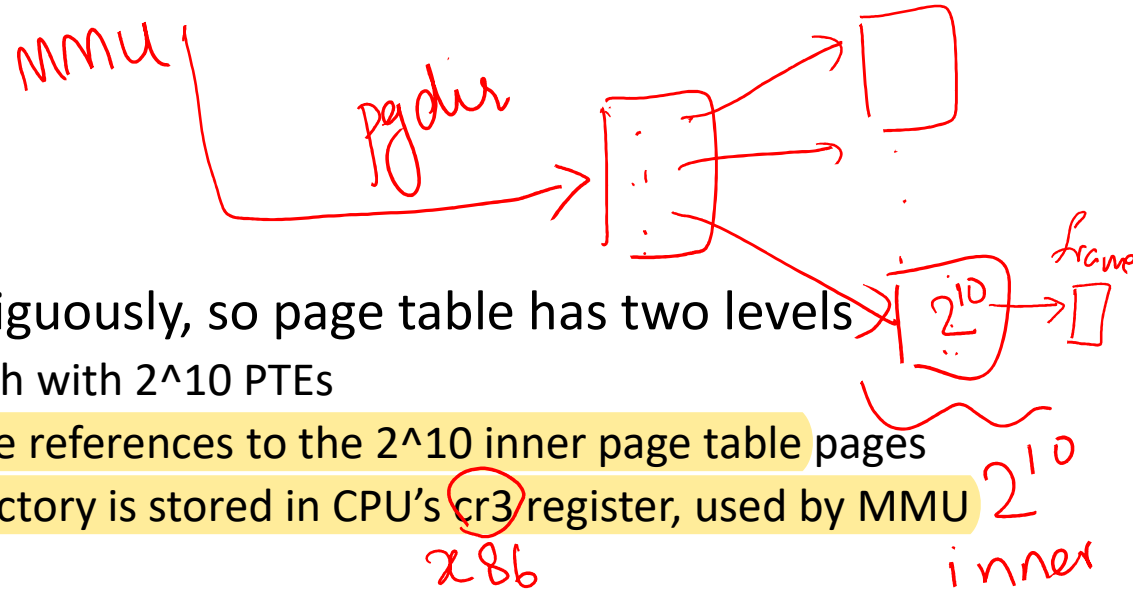
- 32-bit OS, so $2^{32}=4\text{GB}$ virtual address space for every process
- 4KB pages, so 32 bit VA = 20 bit page number + 12 bit offset
- No demand paging, all pages assigned frames always
- Each PTE has 20 bit physical frame number, and some flags
 - PTE_P indicates if page is present (if not set, access will cause page fault)
 - PTE_W indicates if writable (if not set, only reading is permitted)
 - PTE_U indicates if user page (if not set, only kernel can access the page)
- Address translation: use page number (top 20 bits of virtual address) to index into page table, find physical frame number, add 12-bit offset

valid present

20 bit



Two level page table



- 2²⁰ PTEs cannot be stored contiguously, so page table has two levels
 - 2¹⁰ "inner" page table pages, each with 2¹⁰ PTEs
 - Outer page directory stores PTE-like references to the 2¹⁰ inner page table pages
 - Physical address of outer page directory is stored in CPU's cr3 register, used by MMU during address translation
- 32 bit virtual address = 10 bits index into page directory, next 10 bits index into inner page table, last 12 bits are offset within page
 - PFN from PTE + offset = physical address

```

0773 // A virtual address 'la' has a three-part structure as follows:
0774 //
0775 // +-----10-----+-----10-----+-----12-----+
0776 // | Page Directory | Page Table | Offset within Page |
0777 // |   Index      |   Index   |                      |
0778 // +-----+-----+-----+
0779 // \--- PDX(va) ---/ \--- PTX(va) ---/
0780




```

Hand-drawn diagram showing a bracket labeled "MMU" pointing to the virtual address structure.



Process virtual address space in xv6


- Memory image of a process starting at address 0
 - Contains code/data from executable, one page stack (with extra, guard page), expandable heap
- Kernel code/data is mapped beginning at address KERNBASE (2GB)
 - Contains kernel code/data, and free pages maintained by kernel
- Page table of a process contains two sets of PTEs
 - User entries map low virtual addresses to physical memory used by the process for its code/data/stack/heap
 - Kernel entries map high virtual addresses (KERNBASE to KERNBASE+PHYSTOP) to physical memory containing OS code/data and free memory (0 to PHYSTOP)




Maintaining free memory

- After boot up, RAM contains OS code/data and free pages
- OS collects all free pages into a free list, so that it can be assigned to user processes
 - Used for user memory (code/data/stack/heap) and page tables of user processes
- Free list is a linked list, pointer to next free page embedded within previous free page
 - Kernel maintains pointer to first page in the list

```
3115 struct run {  
3116     struct run *next;  
3117 };  
3118  
3119 struct {  
3120     struct spinlock lock;  
3121     int use_lock;  
3122     struct run *freelist;  
3123 } kmem;
```



alloc and free operations




- Anyone who needs a free page calls kalloc()
 - Sets free list pointer to next page and returns first free page on list
- When memory needs to be freed up, kfree() is called
 - Add free page to head of free list, update free list pointer

```

3186 char*
3187 kalloc(void)
3188 {
3189     struct run *r;
3190
3191     if(kmem.use_lock)
3192         acquire(&kmem.lock);
3193     r = kmem.freelist;
3194     if(r)
3195         kmem.freelist = r->next;
3196     if(kmem.use_lock)
3197         release(&kmem.lock);
3198     return (char*)r;
3199 }
  
```

```

3163 void
3164 kfree(char *v)
3165 {
3166     struct run *r;
3167
3168     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
3169         panic("kfree");
3170
3171     // Fill with junk to catch dangling refs.
3172     memset(v, 1, PGSIZE);
3173
3174     if(kmem.use_lock)
3175         acquire(&kmem.lock);
3176     r = (struct run*)v;
3177     r->next = kmem.freelist;
3178     kmem.freelist = r;
3179     if(kmem.use_lock)
3180         release(&kmem.lock);
3181 }
  
```



Memory management of user processes (1)


- User process needs memory pages to build its address space
 - User part of memory image (user code/data/stack/heap)
 - Page table (mappings to user memory image, as well as to kernel code/data)
- Note: only one copy of OS in memory, copied during boot time
- New virtual address space for a process is created during:
 - init process creation
 - fork system call
 - exec system call
- Existing virtual address space modified in sbrk system call
 - Invoked by malloc to expand heap

init fork → Shell

fork
command

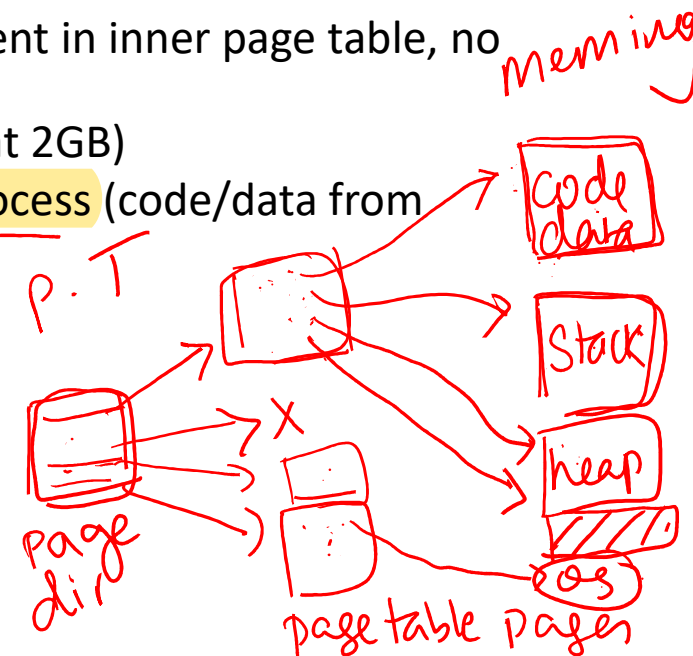
RAM
Contents of process

P.T
of process



Memory management of user processes (2)

- Across all system calls, memory pages obtained from OS via kalloc()
- How is address space of process constructed in fork/exec?
 - Start with one page for the outer page directory
 - Allocate inner page tables on demand (if no entries present in inner page table, no need to allocate a page for it)
 - Add page table mappings for kernel code/data (starting at 2GB)
 - Allocate physical frames to store memory contents of process (code/data from executable, empty stack, ..)
 - Add mappings for user pages in page table
- How is address space of process expanded in sbrk?
 - Allocate physical frames for new virtual addresses
 - Add mappings for newly allocated pages in page table



Recap: fork system call implementation

- Parent allocates new process in ptable, copies parent state to child
- Child process set to runnable, scheduler runs it at a later time
- Return value in parent is PID of child, return value in child is set to 0

→ ret = fork(); →

P → C

```

2579 int
2580 fork(void)
2581 {
2582     int i, pid;
2583     struct proc *np;
2584     struct proc *curproc = myproc();
2585
2586     // Allocate process.
2587     if((np = allocproc()) == 0){
2588         return -1;
2589     }
2590
2591     // Copy process state from proc
2592     if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
2593         kfree(np->kstack);
2594         np->kstack = 0;
2595         np->state = UNUSED;
2596         return -1;
2597     }
2598     np->sz = curproc->sz;
2599     np->parent = curproc;
  
```

kernel
stack

Context
tf:

P Pid
of child

Kernel
mode

memory
copy

```

2600 *np->tf = *curproc->tf;
2601 np->tf->eax = 0;
2602 // Clear %eax so that fork returns 0 in the child.
2603 np->tf->eax = 0;
2604
2605 for(i = 0; i < NOFILE; i++){
2606     if(curproc->ofile[i])
2607         np->ofile[i] = filedup(curproc->ofile[i]);
2608     np->cwd = idup(curproc->cwd);
2609 }
2610 safestrcpy(np->name, curproc->name, sizeof(curproc->name));
2611
2612 pid = np->pid;
2613 acquire(&ptable_lock);
2614
2615 np->state = RUNNABLE;
2616
2617 release(&ptable_lock);
2618
2619 return pid;
2620 }
2621
  
```

copy trapframe

} files

np

ptable

Fork: copying memory image

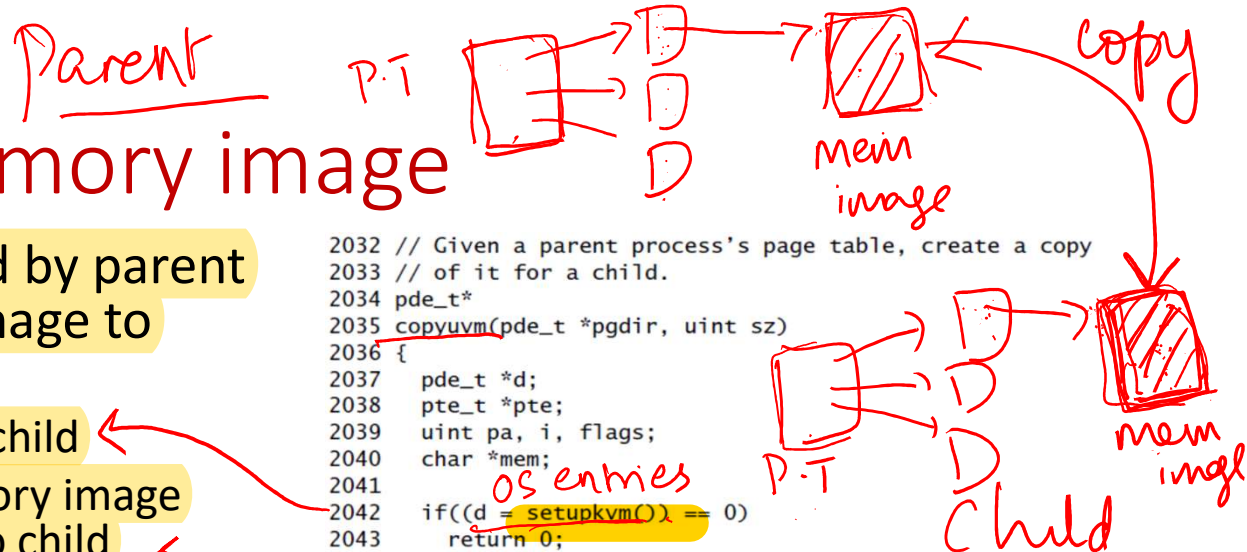
- Function “copyuvm” called by parent to copy parent memory image to child

- Create new page table for child
- Walk through parent memory image page by page and copy it to child

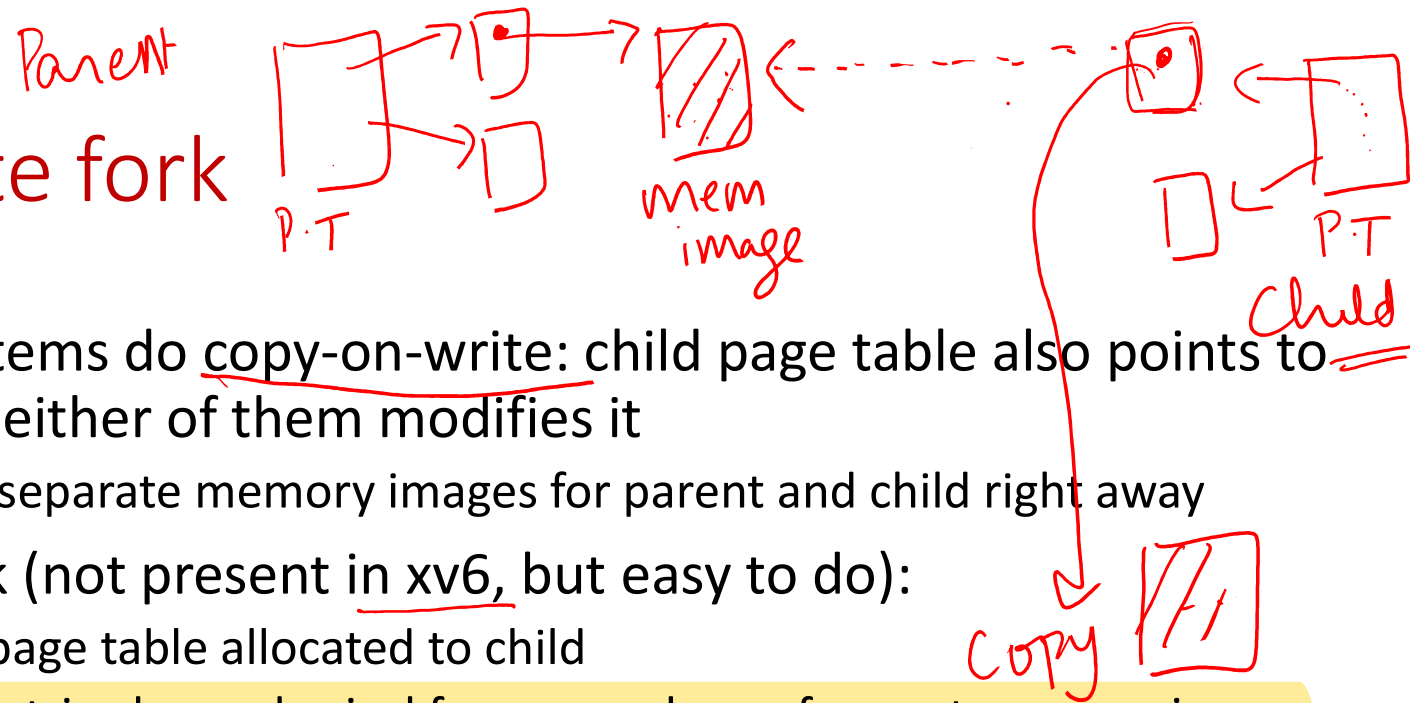
- For each page in parent
 - Fetch PTE, get physical address, permissions
 - Allocate new frame for child, copy contents of parent's page to new page of child
 - Add a PTE from virtual address to physical address of new page in child page table

```

2032 // Given a parent process's page table, create a copy
2033 // of it for a child.
2034 pde_t*
2035 copyuvm(pde_t *pgdir, uint sz)
2036 {
2037     pde_t *d;
2038     pte_t *pte;
2039     uint pa, i, flags;
2040     char *mem;
2041
2042     if((d = setupkvm()) == 0)
2043         return 0;
2044     for(i = 0; i < sz; i += PGSIZE){
2045         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
2046             panic("copyuvm: pte should exist");
2047         if(!(*pte & PTE_P))
2048             panic("copyuvm: page not present");
2049         pa = PTE_ADDR(*pte);
2050         flags = PTE_FLAGS(*pte);
2051         if((mem = kalloc()) == 0)
2052             goto bad;
2053         memmove(mem, (char*)P2V(pa), PGSIZE);
2054         if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
2055             kfree(mem);
2056             goto bad;
2057         }
2058     }
2059     return d;
2060
2061 bad:
2062     freevm(d);
2063     return 0;
2064 }
    
```



Copy-on-write fork




- Real operating systems do copy-on-write: child page table also points to parent pages until either of them modifies it
 - Here, xv6 creates separate memory images for parent and child right away
- Copy-on-write fork (not present in xv6, but easy to do):
 - During fork, new page table allocated to child
 - Child page table entries have physical frame numbers of parent memory image pages only, no copy created for child
 - Parent's memory image is marked as read only
 - When parent or child tries to modify, MMU traps to OS
 - As part of trap handling, separate copy of memory image created
 - Finally, two separate copies of memory image for parent and child

Growing memory image: sbrk

- Initially heap is empty, program “break” (end of user memory) is at end of stack
 - sbrk() system call invoked by malloc to expand heap
- To grow memory, allocuvm allocates new pages, adds mappings into page table for new pages
- Whenever page table updated, must update cr3 register and TLB (done even during context switching)

```
2557 int
2558 growproc(int n)
2559 {
2560     uint sz;
2561     struct proc *curproc = myproc();
2562
2563     sz = curproc->sz;
2564     if(n > 0){
2565         if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
2566             return -1;
2567     } else if(n < 0){
2568         if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
2569             return -1;
2570     }
2571     curproc->sz = sz;
2572     switchuvm(curproc);
2573     return 0;
2574 }
```

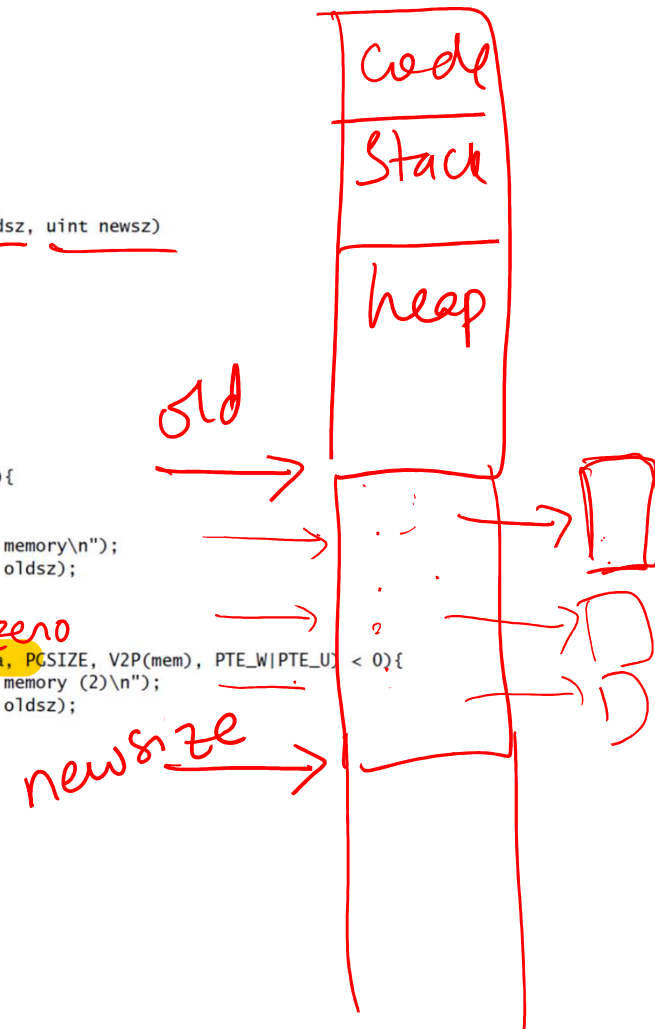


← update pagetable
MMU

allocuvvm: grow address space

- Walk through new virtual addresses, page by page
- Allocate new frame, add mapping to page table with suitable user permissions
- Similarly deallocuvvm shrinks memory image, frees up pages

```
1926 int
1927 allocuvvm(pde_t *pgdir, uint oldsz, uint newsz)
1928 {
1929     char *mem;
1930     uint a;
1931
1932     if(newsz >= KERNBASE)
1933         return 0;
1934     if(newsz < oldsz)
1935         return oldsz;
1936
1937     a = PGROUNDUP(oldsz);
1938     for(; a < newsz; a += PGSIZE){
1939         mem = kalloc();
1940         if(mem == 0){
1941             cprintf("allocuvvm out of memory\n");
1942             deallocuvvm(pgdir, newsz, oldsz);
1943             return 0;
1944         }
1945         memset(mem, 0, PGSIZE);
1946         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
1947             cprintf("allocuvvm out of memory (2)\n");
1948             deallocuvvm(pgdir, newsz, oldsz);
1949             kfree(mem);
1950             return 0;
1951         }
1952     }
1953     return newsz;
1954 }
```



Exec system call (1)

- Read ELF binary file from disk into memory
- Start with new page table, add mappings to new executable pages and grow virtual address space
- Do not overwrite old page table yet

```

6609 int
6610 exec(char *path, char **argv)
6611 {
6612     char *s, *last;
6613     int i, off;
6614     uint argc, sz, sp, ustack[3+MAXARG+1];
6615     struct elfhdr elf;
6616     struct inode *ip;
6617     struct proghdr ph;
6618     pde_t *pgdir, *oldpgdir;
6619     struct proc *curproc = myproc();
6620
6621     begin_op();
6622
6623     if((ip = namei(path)) == 0){
6624         end_op();
6625         cprintf("exec: fail\n");
6626         return -1;
6627     }
6628     ilock(ip);
6629     pgdir = 0;
6630
6631     // Check ELF header
6632     if(readi(ip, (char*)&elf, 0, sizeof(elf)) != sizeof(elf))
6633         goto bad;
6634     if(elf.magic != ELF_MAGIC)
6635         goto bad;
6636
6637     if((pgdir = setupkvm()) == 0)
6638         goto bad;

```



read from
disk into
memory
& add PTE
Start new pgdir

```

6640 // Load program into memory.
6641 sz = 0;
6642 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
6643     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
6644         goto bad;
6645     if(ph.type != ELF_PROG_LOAD)
6646         continue;
6647     if(ph.memsz < ph.filesz)
6648         goto bad;
6649     if(ph.vaddr + ph.memsz < ph.vaddr)
6650         goto bad;
6651     if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
6652         goto bad;
6653     if(ph.vaddr % PGSIZE != 0)
6654         goto bad;
6655     if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
6656         goto bad;
6657 }
6658 iunlockput(ip);
6659 end_op();
6660 ip = 0;

```


Copy
new
exec
from disk



Exec system call (2)

`exec(cmd, " ", " ")`

a.out



- After executable is copied to memory image, allocate 2 pages for stack (one is guard page, permissions cleared, access will trap)
- Push exec arguments onto user stack for main function of new program

```
6662 // Allocate two pages at the next page boundary.
6663 // Make the first inaccessible. Use the second as the user stack.
6664 sz = PGROUNDUP(sz);
6665 if((sz = allocvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
6666     goto bad;
6667 clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
6668 sp = sz;
6669
6670 // Push argument strings, prepare rest of stack in ustack.
6671 for(argc = 0; argv[argc]; argc++) {
6672     if(argc >= MAXARG)
6673         goto bad;
6674     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
6675     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
6676         goto bad;
6677     ustack[3+argc] = sp;
6678 }
6679 ustack[3+argc] = 0;
6680
6681 ustack[0] = 0xffffffff; // fake return PC
6682 ustack[1] = argc;
6683 ustack[2] = sp - (argc+1)*4; // argv pointer
6684
6685 sp -= (3+argc+1) * 4;
6686 if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
6687     goto bad;
6688
```

2 pages stack

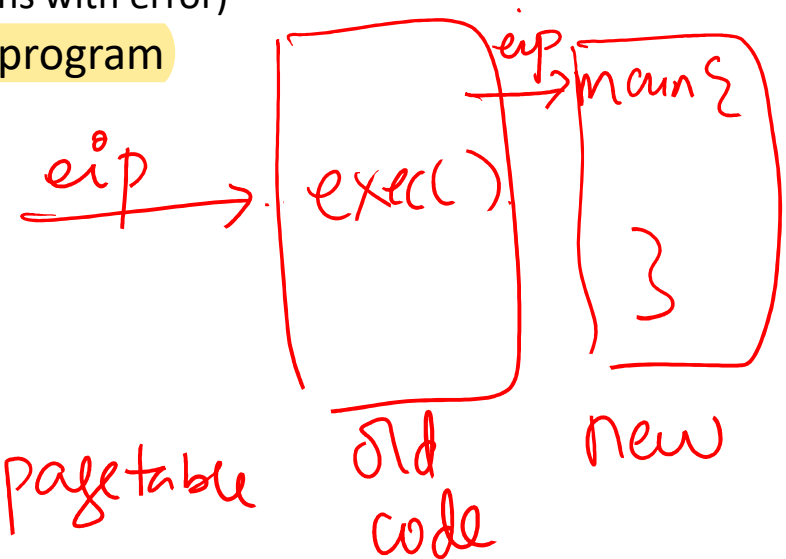
command line arguments

main(argc, argv)
{
 local ...
}

Exec system call (3)

- If no errors so far, switch to new page table that is pointing to new memory image
 - If any error, go back to old memory image (exec returns with error)
- Set eip in trapframe to start at entry point of new program
 - Returning from trap, process will run new executable

```
6689 // Save program name for debugging.
6690 for(last=s=path; *s; s++)
6691     if(*s == '/')
6692         last = s+1;
6693 safestrcpy(curproc->name, last, sizeof(curproc->name));
6694
6695 // Commit to the user image.
6696 oldpgdir = curproc->pgdir;
6697 curproc->pgdir = pgdir;
6698 curproc->sz = sz;
6699 curproc->tf->eip = elf.entry; // main
6700 curproc->tf->esp = sp;
6701 switchvm(curproc);
6702 freevm(oldpgdir);
6703 return 0;
6704
6705 bad:
6706 if(pgdir)
6707     freevm(pgdir);
6708 if(ip){
6709     iunlockput(ip);
6710     end_op();
6711 }
6712 return -1;
6713 }
```



Switch to new pagetable

free up old memory