

CS 347M (Operating Systems Minor)

Spring 2022

# Lecture 12: Threads and Concurrency

Mythili Vutukuru  
CSE, IIT Bombay


# Recap: Growing memory image: sbrk

- Initially heap is empty, program “break” (end of user memory) is at end of stack
  - sbrk() system call invoked by malloc to expand heap
- To grow memory, allocuvm allocates new pages, adds mappings into page table
- Whenever page table updated, must update cr3 register and TLB (done even during context switching)

```
2557 int
2558 growproc(int n)
2559 {
2560     uint sz;
2561     struct proc *curproc = myproc();
2562
2563     sz = curproc->sz;
2564     if(n > 0){
2565         if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
2566             return -1;
2567     } else if(n < 0){
2568         if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
2569             return -1;
2570     }
2571     curproc->sz = sz;
2572     switchuvm(curproc);
2573     return 0;
2574 }
```

program break

heap



# allocvm: grow address space

- Walk through new virtual addresses to be added in page size chunks
- Allocate new page, add it to page table with suitable user permissions
- Similarly deallocvm shrinks memory image, frees up pages

```
1926 int
1927 allocvm(pde_t *pgdir, uint oldsz, uint newsz)
1928 {
1929     char *mem;
1930     uint a;
1931
1932     if(newsz >= KERNBASE)
1933         return 0;
1934     if(newsz < oldsz)
1935         return oldsz;
1936
1937     a = PGROUNDUP(oldsz);
1938     for(; a < newsz; a += PGSIZE){
1939         mem = kalloc();
1940         if(mem == 0){
1941             cprintf("allocvm out of memory\n");
1942             deallocvm(pgdir, newsz, oldsz);
1943             return 0;
1944         }
1945         memset(mem, 0, PGSIZE);
1946         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
1947             cprintf("allocvm out of memory (2)\n");
1948             deallocvm(pgdir, newsz, oldsz);
1949             kfree(mem);
1950             return 0;
1951         }
1952     }
1953     return newsz;
1954 }
```


*add PTE*

# Functions to build/walk page table

- Page table entries added by “mappages”
  - Arguments: page directory, range of virtual addresses, physical addresses to map to, permissions of the pages
  - For each page, walks page table, get pointer to PTE via function “walkpgdir”, fills it with physical addr, permissions
- Function “walkpgdir” walks page table, returns PTE of a virtual address
  - Can allocate inner page table if it doesn't exist (based on value of last argument)



```
1756 // Create PTEs for virtual addresses starting at va that refer to
1757 // physical addresses starting at pa. va and size might not
1758 // be page-aligned.
1759 static int
1760 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
1761 {
1762     char *a, *last;
1763     pte_t *pte;
1764
1765     a = (char*)PGROUNDDOWN((uint)va);
1766     last = (char*)PGROUNDDOWN((uint)va) + size - 1;
1767     for(;;){
1768         if((pte = walkpgdir(pgdir, a, 1)) == 0) find PTE
1769             return -1;
1770         if(*pte & PTE_P)
1771             panic("remap");
1772         *pte = pa | perm | PTE_P; add PA to PTE
1773         if(a == last)
1774             break;
1775         a += PGSIZE;
1776         pa += PGSIZE;
1777     }
1778     return 0;
1779 }
```

```
1731 // Return the address of the PTE in page table pgdir
1732 // that corresponds to virtual address va. If alloc!=0,
1733 // create any required page table pages.
1734 static pte_t *
1735 walkpgdir(pde_t *pgdir, const void *va, int alloc)
1736 {
1737     pde_t *pde;
1738     pte_t *pgtab;
1739
1740     pde = &pgdir[PDX(va)];
1741     if(*pde & PTE_P){
1742         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
1743     } else {
1744         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
1745             return 0;
1746         // Make sure all those PTE_P bits are zero.
1747         memset(pgtab, 0, PGSIZE);
1748         // The permissions here are overly generous, but they can
1749         // be further restricted by the permissions in the page table
1750         // entries, if necessary.
1751         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
1752     }
1753     return &pgtab[PTX(va)];
1754 }
```



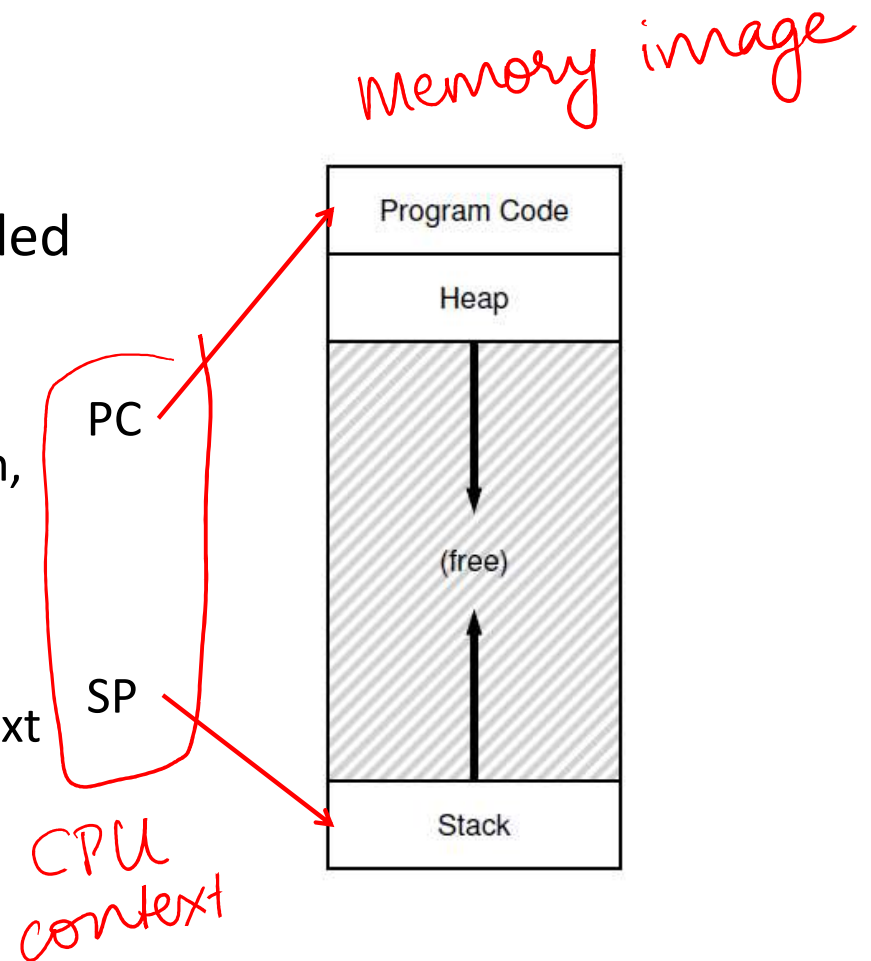
# Programming assignment 2

- New system calls in xv6:
  - Calculate number of pages in virtual address space
    - Process size in PCB / PGSIZE
  - Calculate number of physical pages given to process
    - Walk page table, see how many pages have present flag set
- New syscall: memory map a page and allocate memory on demand
  - Much like sbrk, add pages at program break
  - Currently sbrk system call increase both virtual and physical memory
  - Your new mmap syscall should only increase virtual memory
  - Physical memory should be allocated when page accessed and page fault
  - Page fault handler should allocate physical frame



# Processes and threads

- So, far we have studied single threaded programs
- Recap: process execution
  - CPU executes instruction by instruction, traps to OS as needed
  - PC points to next instruction to run
  - SP points to current top of stack
  - Other registers also with process context
- A program can also have multiple threads of execution
- What is a thread?



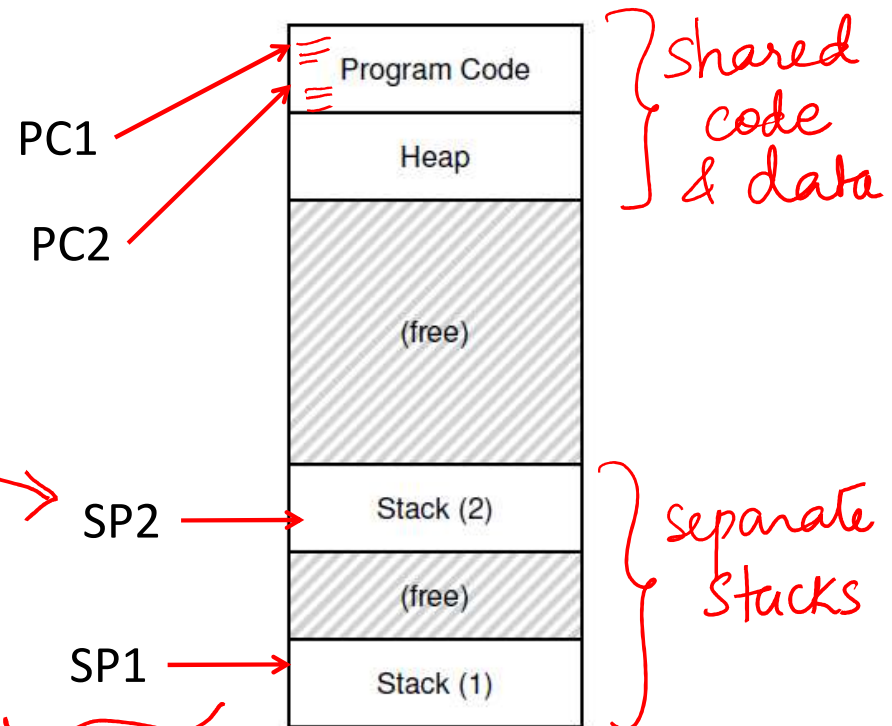
# What are threads?

- Threads = light weight processes
- Why? A process may want to run multiple copies of itself
  - If one copy blocks due to blocking system call, another copy can still run
  - Multiple copies can run in parallel on multiple CPU cores
- Why not have multiple child processes running the same program?
  - Disadvantage: too much memory consumed by identical memory images
- A process can create multiple threads (default: single thread)
  - Multiple threads share same memory image of process, saves memory
  - Threads run independently on same code (if one blocks, another can still run)
  - Threads can run in parallel on multiple cores at same time



# Multi-threaded process

- A thread is like another copy of a process that executes independently from parent
- Threads shares the same code, global/static data, heap
- Each thread has separate stack for independent function calls
- Each thread has separate PC
  - Each thread may run over different part of the program independently
- Each thread has separate CPU context during execution





# Concurrency vs. parallelism




- Understand the difference between concurrency and parallelism
  - **Concurrency:** running multiple threads/processes at the same time, even on single CPU core, by interleaving their executions
  - **Parallelism:** running multiple threads/processes in parallel over different CPU cores
- With multiple threads, process can get better performance on multicore systems via parallelism
- Even if no parallelism (single core), concurrency of threads ensures effective use of CPU when one of the threads blocks (e.g., for I/O)

# POSIX threads

- In Linux, POSIX threads (**pthread**s) library allows creation of multiple threads in a process
- Each thread is given a **start function** where its execution begins
  - Threads execute independently from parent after creation
  - Parent can wait for threads to finish (optional)
- Several such threading libraries exist in different programming languages

```
void f1() {  
    ...  
}  
  
void f2() {  
    ...  
}  
  
main() {  
    ...  
    pthread_t t1, t2  
    pthread_create(&t1, .., f1,..)  
    pthread_create(&t2, .., f2,..)  
    ...  
  
    pthread_join(t1, ..)  
    pthread_join(t2, ..)  
  
}
```



# Creating threads using pthreads API

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>
4
5  void *mythread(void *arg) {
6      printf("%s\n", (char *) arg);
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     pthread_t p1, p2;
13     int rc;
14     printf("main: begin\n");
15     rc = pthread_create(&p1, NULL, mythread, "A"); assert(rc == 0);
16     rc = pthread_create(&p2, NULL, mythread, "B"); assert(rc == 0);
17     // join waits for the threads to finish
18     rc = pthread_join(p1, NULL); assert(rc == 0);
19     rc = pthread_join(p2, NULL); assert(rc == 0);
20     printf("main: end\n");
21     return 0;
22 }
```

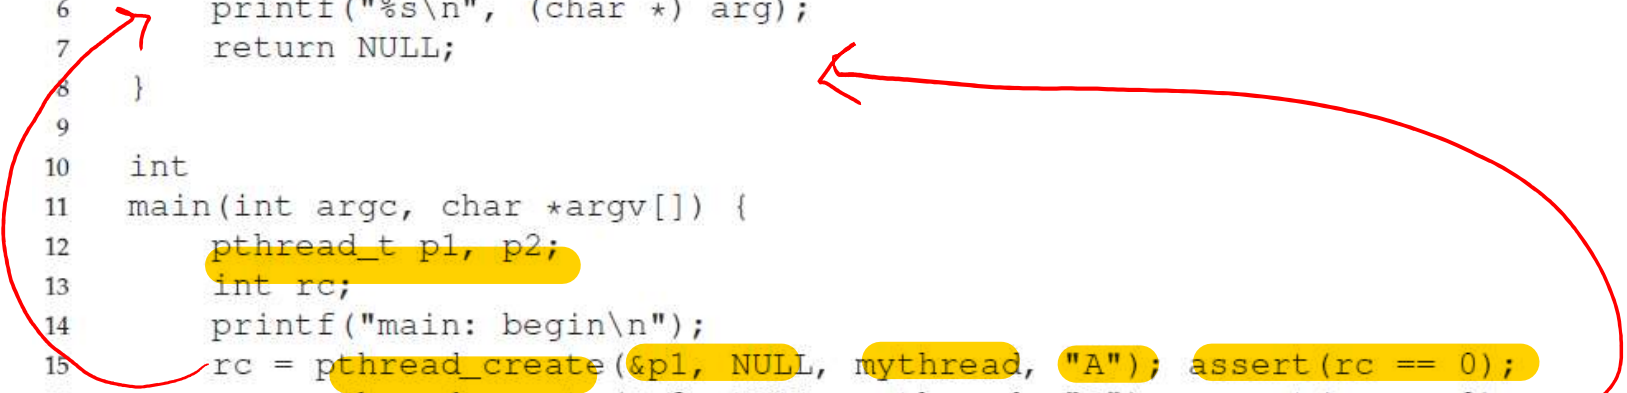


Figure 26.2: Simple Thread Creation Code (t0.c)

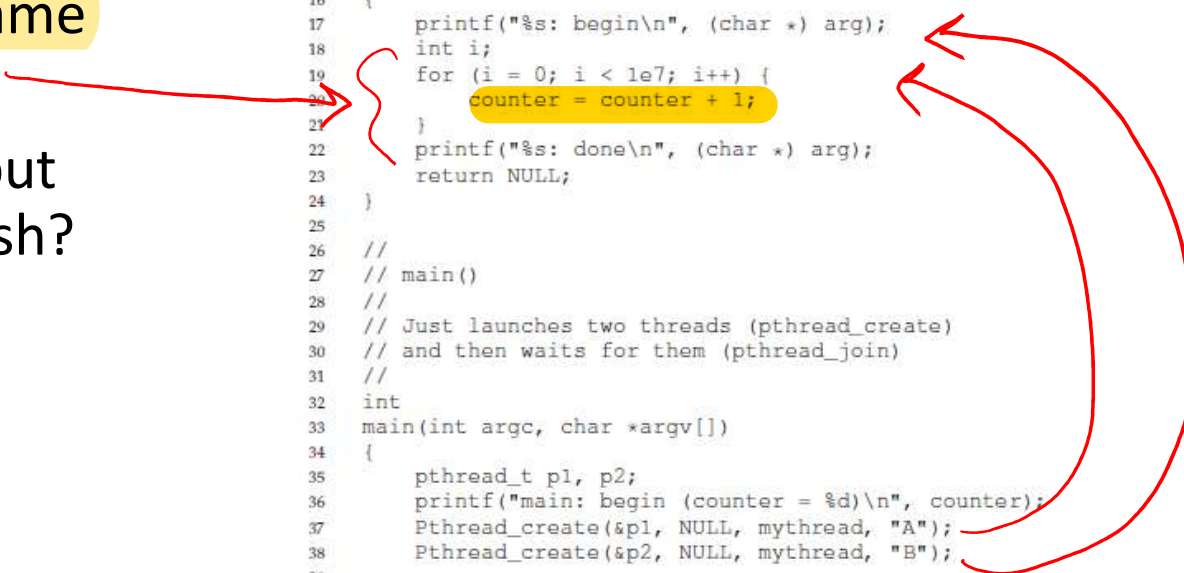
# Scheduling threads

- OS schedules threads that are ready to run independently, like processes
- The context of a thread (PC, registers) is saved into/restored from thread control block (TCB)
  - Every PCB has one or more linked TCBs
- Threads that are scheduled independently by kernel are called kernel threads
  - E.g., Linux pthreads are kernel threads
- In contrast, some libraries provide user-level threads
  - User program sees multiple threads, but kernel is aware of fewer threads
  - Multiple such user threads may not be scheduled in parallel by kernel
  - Why use user threads then? Ease of programming

# Example: threads with shared data


- Shared global counter
- Two threads update same counter  $10^7$  times
- What is expected output after both threads finish?

```
4
5 static volatile int counter = 0;
6
7 //
8 // mythread()
9 //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *
15 mythread(void *arg)
16 {
17     printf("%s: begin\n", (char *) arg);
18     int i;
19     for (i = 0; i < 1e7; i++) {
20         counter = counter + 1;
21     }
22     printf("%s: done\n", (char *) arg);
23     return NULL;
24 }
25
26 //
27 // main()
28 //
29 // Just launches two threads (pthread_create)
30 // and then waits for them (pthread_join)
31 //
32 int
33 main(int argc, char *argv[])
34 {
35     pthread_t p1, p2;
36     printf("main: begin (counter = %d)\n", counter);
37     Pthread_create(&p1, NULL, mythread, "A");
38     Pthread_create(&p2, NULL, mythread, "B");
39
40     // join waits for the threads to finish
41     Pthread_join(p1, NULL);
42     Pthread_join(p2, NULL);
43     printf("main: done with both (counter = %d)\n", counter);
44     return 0;
45 }
```

Two red arrows originate from the list on the left. One arrow points from the text 'Two threads update same counter 10^7 times' to line 20 of the code, which contains 'counter = counter + 1;'. The other arrow points from the same list item to line 37 of the code, which contains 'Pthread\_create(&p1, NULL, mythread, "A");'. A third red arrow points from the text 'What is expected output after both threads finish?' to line 43 of the code, which contains 'printf("main: done with both (counter = %d)\n", counter);'.

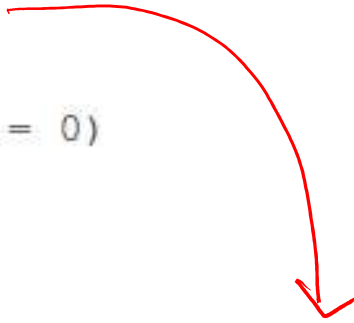
## Threads with shared data: what happens?

- What do we expect? Two threads, each increments counter by  $10^7$ , so  $2 \times 10^7$



```
prompt> gcc -o main main.c -Wall -pthread
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 20000000)
```

- Sometimes, a lower value. Why?



```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)
```

# Understanding shared data access

load counter → reg  
reg = reg + 1  
store reg → counter

- The C code “counter = counter + 1” is compiled into multiple instructions
  - Load counter variable from memory into register
  - Increment register
  - Store register back into memory of counter variable
- What happens when two threads run this line of code concurrently?

- Counter is 0 initially
- T1 loads counter into register, increment reg
- Context switch, register (value 1) saved
- T2 runs, loads counter 0 from memory
- T2 increments register, stores to memory
- T1 resumes, stores register value to counter
- Counter value rewritten to 1 again
- Final counter value is 1, expected value is 2

T1

load counter → reg  
reg = reg + 1  
**(context switch, save reg)**

T2

load counter → reg  
reg = reg + 1  
store reg → counter

**(resume, restore reg)**  
store reg → counter

# Race conditions, critical sections

- Incorrect execution of code due to concurrency is called **race condition**
  - Due to **unfortunate timing of context switches**, atomicity of data update violated
- Race conditions happen when we have **concurrent execution on shared data**
  - **Threads** sharing common data in memory image of user processes
  - Processes in kernel mode sharing **OS data structures**
- We require **mutual exclusion** on some parts of user or OS code
  - Concurrent execution by multiple threads/processes should not be permitted
- Parts of program that need to be executed with mutual exclusion for correct operation are called **critical sections**
  - Present in **multi-threaded programs, OS code**
- How to **access critical sections with mutual exclusion?** Using locks (next topic)