

CS 347M (Operating Systems Minor)

Spring 2022

Lecture 15: Sleep and wakeup in xv6

Mythili Vutukuru
CSE, IIT Bombay

Locks, sleep/wakeup in xv6

- xv6 does not have userspace threads, only single threaded processes
- But multiple processes may be kernel mode on different CPU
 - Need **locks** to protect access to shared kernel data structures
- OS also needs a mechanism to let processes **sleep** (e.g., when process makes blocking disk read syscall) and **wakeup** when some events occur (e.g., disk has raised interrupt and data is ready)
 - Needs sleep/wakeup functions for processes in kernel mode (not userspace)
 - Process P1 in kernel mode calls sleep to give up CPU, gets blocked until event
 - Another process P2 (in kernel mode) wakes up P1 when the event occurs
- This lecture: more on xv6 locks, sleep, wakeup, ...

Recap: Context switching in xv6 (1)

- Every CPU has a scheduler thread (special process that runs scheduler code)
- Scheduler goes over list of processes and switches to one of the runnable ones
- The special function “swtch” performs the actual context switch
 - Save context on kernel stack of old process
 - Restore context from kernel stack of new process

```
2757 void
2758 scheduler(void)
2759 {
2760     struct proc *p;
2761     struct cpu *c = mycpu();
2762     c->proc = 0;
2763
2764     for(;;){
2765         // Enable interrupts on this processor.
2766         sti();
2767
2768         // Loop over process table looking for process to run.
2769         acquire(&ptable.lock);
2770         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2771             if(p->state != RUNNABLE)
2772                 continue;
2773
2774             // Switch to chosen process. It is the process's job
2775             // to release ptable.lock and then reacquire it
2776             // before jumping back to us.
2777             c->proc = p;
2778             switchvm(p);
2779             p->state = RUNNING;
2780
2781             swtch(&c->scheduler, p->context);
2782             switchkvm();
2783
2784             // Process is done running for now.
2785             // It should have changed its p->state before coming back.
2786             c->proc = 0;
2787         }
2788         release(&ptable.lock);
2789     }
2790 }
2791 }
```

Recap: Context switching in xv6 (2)

- After running for some time, the process switches back to the scheduler thread, when:
 - Process has terminated (exit system call)
 - Process needs to sleep (e.g., blocking read system call)
 - Process yields after running for long (timer interrupt)
- Process calls “sched” which calls “swtch” to switch to scheduler thread again
- Scheduler thread runs its loop and picks next process to run, and the story repeats

```
2662 // Jump into the scheduler, never to return.
2663 curproc->state = ZOMBIE;
2664 sched();
2665 panic("zombie exit");
2666 }
```


```
2894 // Go to sleep.
2895 p->chan = chan;
2896 p->state = SLEEPING;
2897
2898 sched();
2899
```

```
2826 // Give up the CPU for one scheduling round.
2827 void
2828 yield(void)
2829 {
2830     acquire(&ptable.lock);
2831     myproc()->state = RUNNABLE;
2832     sched();
2833     release(&ptable.lock);
2834 }
```

ptable.lock (1)

```
2409 struct {  
2410     struct spinlock lock;  
2411     struct proc proc[NPROC];  
2412 } ptable;
```

- The process table protected by a lock, any access to ptable must be done with ptable.lock held
- Normally, a process in kernel mode acquires ptable.lock, changes ptable in some way, releases lock
 - Example: when allocproc allocates new struct proc
- But during context switch from process P1 to P2, ptable structure is being changed all through context switch, so when to release lock?
 - P1 acquires lock, switches to scheduler, switches to P2, P2 releases lock




ptable.lock (2)

- Every function that calls sched() to give up CPU will do so with ptable.lock held
- Which functions invoke sched() to give up CPU?
 - Yield: process gives up CPU due to timer interrupt
 - Sleep: when process wishes to block
 - Exit: when process terminates
- Every function where a process resumes after being scheduled release ptable.lock
- What functions does a process resume after swtch?
 - Yield: resuming process after yield is done
 - Sleep: resuming process that is waking up after sleep
 - Forkret: for newly created processes
- Purpose of forkret: to release ptable.lock
 - New process then returns from trap like its parent

```
2826 // Give up the CPU for one scheduling round.
2827 void
2828 yield(void)
2829 {
2830     acquire(&ptable.lock);
2831     myproc()->state = RUNNABLE;
2832     sched();
2833     release(&ptable.lock);
2834 }
```

```
2852 void
2853 forkret(void)
2854 {
2855     static int first = 1;
2856     // Still holding ptable.lock from scheduler.
2857     release(&ptable.lock);
2858
2859     if (first) {
2860         // Some initialization functions must be run i
2861         // of a regular process (e.g., they call sleep
2862         // be run from main().
2863         first = 0;
2864         iinit(ROOTDEV);
2865         initlog(ROOTDEV);
2866     }
```

ptable.lock (3)



- Scheduler goes into loop with lock held
- Acquire ptable.lock in P1 \rightarrow scheduler picks P2 \rightarrow release in P2
- Later, acquire ptable.lock in P2 \rightarrow scheduler picks P3 \rightarrow release in P3
- Periodically, end of looping over all processes, releases lock temporarily
 - What if no runnable process found due to interrupts being disabled? Release lock, enable interrupts, allow processes to become runnable.

```
2757 void
2758 scheduler(void)
2759 {
2760     struct proc *p;
2761     struct cpu *c = mycpu();
2762     c->proc = 0;
2763
2764     for(;;){
2765         // Enable interrupts on this processor.
2766         sti();
2767
2768         // Loop over process table looking for process to run.
2769         acquire(&ptable.lock);
2770         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
2771             if(p->state != RUNNABLE)
2772                 continue;
2773
2774             // Switch to chosen process. It is the process's job
2775             // to release ptable.lock and then reacquire it
2776             // before jumping back to us.
2777             c->proc = p;
2778             switchvm(p);
2779             p->state = RUNNING;
2780
2781             swtch(&(c->scheduler), p->context);
2782             switchkvm();
2783
2784             // Process is done running for now.
2785             // It should have changed its p->state before coming back.
2786             c->proc = 0;
2787         }
2788         release(&ptable.lock);
2789     }
2790 }
2791 }
```

Sleep and wakeup in xv6

- A process P1 that wishes to block and give up CPU calls “sleep” function
 - Example: process reads a block from disk, must block until disk read completes
 - Read syscall → sleep → sched() to give up CPU
- Another process P2 calls “wakeup” when event to block P1 occurs
 - P2 calls wakeup → marks P1 as runnable, no context switch immediately
 - Example: disk interrupt occurred when P2 is running, P2 runs interrupt handler, which will call wakeup
- How does P2 know which process to wake up? When P1 sleeps, it sets a channel (void * chan) in its struct proc, P2 calls wakeup on same channel
 - Channel = any value known to both P1 and P2
 - Example: channel value for disk read can be address of disk block
- Spinlock protects atomicity of sleep: P1 calls sleep with some spinlock L held, P2 calls wakeup with same spinlock L held

Sleep function

- Two arguments: channel to sleep on, a spinlock to protect atomicity of sleeping
- Acquire ptable.lock, release the lock given to sleep (make it available for wakeup)
 - Unless lock given is ptable.lock itself, in which case no need to acquire again
 - One of two locks held at all times
- Sleep calls sched() to give up CPU
 - Needs to hold ptable.lock when calling sched()
- Calls sched(), switched out of CPU, resumes again when woken up and ready to run
- Reacquires the lock given to sleep and returns back
 - Code that invoked sleep with lock held returns with lock held again

```
2871 // Atomically release lock and sleep on chan.
2872 // Reacquires lock when awakened.
2873 void
2874 sleep(void *chan, struct spinlock *lk)
2875 {
2876     struct proc *p = myproc();
2877
2878     if(p == 0)
2879         panic("sleep");
2880
2881     if(lk == 0)
2882         panic("sleep without lk");
2883
2884     // Must acquire ptable.lock in order to
2885     // change p->state and then call sched.
2886     // Once we hold ptable.lock, we can be
2887     // guaranteed that we won't miss any wakeup
2888     // (wakeup runs with ptable.lock locked),
2889     // so it's okay to release lk.
2890     if(lk != &ptable.lock){
2891         acquire(&ptable.lock);
2892         release(lk);
2893     }
2894     // Go to sleep.
2895     p->chan = chan;
2896     p->state = SLEEPING;
2897
2898     sched();
2899
2900     // Tidy up.
2901     p->chan = 0;
2902
2903     // Reacquire original lock.
2904     if(lk != &ptable.lock){
2905         release(&ptable.lock);
2906         acquire(lk);
2907     }
2908 }
```

Wakeup function


- Sleep and wakeup called by processes with same lock held (protect atomicity of sleep)
- Wakeup acquires ptable.lock, changes ptable to mark process matching channel as runnable, releases ptable.lock
 - If lock protecting atomicity of sleep is ptable.lock itself, then directly call wakeup1
- Sleep holds one of sleep's lock or ptable.lock at all times, so a wakeup cannot run in between sleep
- Wakes up all processes sleeping on a channel in ptable (more like signal broadcast of condition variables)
 - Good idea to check condition is still true upon waking up (use while loop while calling sleep)

```
2950 // Wake up all processes sleeping on chan.
2951 // The ptable lock must be held.
2952 static void
2953 wakeup1(void *chan)
2954 {
2955     struct proc *p;
2956
2957     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
2958         if(p->state == SLEEPING && p->chan == chan)
2959             p->state = RUNNABLE;
2960 }
2961
2962 // Wake up all processes sleeping on chan.
2963 void
2964 wakeup(void *chan)
2965 {
2966     acquire(&ptable.lock);
2967     wakeup1(chan);
2968     release(&ptable.lock);
2969 }
```

Example: pipes (1)

- xv6 provides anonymous pipes for IPC between parent and child processes
 - E.g., Parent P and child C share anonymous pipe
 - Child C writes into pipe, parent P reads from pipe
- Internal implementation inside kernel
 - Common shared buffer, protected by a spinlock
 - Write system call stores data in shared buffer
 - Read system call returns data from shared buffer
- Sleep and wakeup involved in read/write
 - Pipe read sleeps if pipe is empty, pipe write wake up
 - Pipe write sleeps if pipe is full, pipe read wakes up

```
6762 struct pipe {
6763     struct spinlock lock;
6764     char data[PIPESIZE];
6765     uint nread;    // number of bytes read
6766     uint nwrite;   // number of bytes written
6767     int readopen;  // read fd is still open
6768     int writeopen; // write fd is still open
6769 };
```



```
//userspace code

int fd[2]
pipe(fd) //syscall to create pipe

int ret = fork()

if(ret == 0) { //child
    close(fd[0]) //close read end
    write(fd[1], message, ..)
}
else { //parent
    close(fd[1]) //close write end
    read(fd[0], message, ..)
}
```

Example: pipes (2)

- Implementation of pipe read and write system calls using sleep/wakeup
 - Similar to producer-consumer logic of last class
 - Channel for sleep/wakeup = address of pipe structure variables (can be

```
6829 int
6830 pipewrite(struct pipe *p, char *addr, int n)
6831 {
6832     int i;
6833
6834     acquire(&p->lock);
6835     for(i = 0; i < n; i++){
6836         while(p->nwrite == p->nread + PIPESIZE){ pipe is full
6837             if(p->readopen == 0 || myproc()->killed){
6838                 release(&p->lock);
6839                 return -1;
6840             }
6841             wakeup(&p->nread);
6842             sleep(&p->nwrite, &p->lock); writer's channel for sleep is
6843         } address of nwrite variable
6844         p->data[p->nwrite++ % PIPESIZE] = addr[i];
6845     }
6846     wakeup(&p->nread);
6847     release(&p->lock);
6848     return n;
6849 }
```

```
6850 int
6851 piperead(struct pipe *p, char *addr, int n)
6852 {
6853     int i;
6854
6855     acquire(&p->lock);
6856     while(p->nread == p->nwrite && p->writeopen){
6857         if(myproc()->killed){ pipe is empty
6858             release(&p->lock);
6859             return -1;
6860         }
6861         sleep(&p->nread, &p->lock); pipe lock protects
6862     } atomicity of sleep
6863     for(i = 0; i < n; i++){
6864         if(p->nread == p->nwrite)
6865             break;
6866         addr[i] = p->data[p->nread++ % PIPESIZE];
6867     }
6868     wakeup(&p->nwrite);
6869     release(&p->lock);
6870     return i;
6871 }
```

Example: wait and exit

- If wait called in parent while children are running, parent calls sleep and gives up CPU
 - Here, channel is parent struct proc address, lock given to sleep is ptable.lock

```
2706    // Wait for children to exit. (See wakeup1 call in proc_exit.)
2707    sleep(curproc, &ptable.lock);
```

- In exit, child acquires ptable.lock and wakes up sleeping parent using its channel

```
2650    // Parent might be sleeping in wait().
2651    wakeup1(curproc->parent);
```

- Here, lock given to protect atomicity of sleep is ptable.lock itself (convenient to do so)
 - Double locking of ptable.lock is avoided during sleep and wakeup
- Why is terminated process memory cleaned up by parent?
 - When a process calls exit, kernel stack, page table etc are in use, all this memory cannot be cleared until terminated process has been taken off the CPU
 - Parent code in wait is a good place to clean up child memory after child has stopped running