

CS 347M (Operating Systems Minor)

Spring 2022

Lecture 16: Semaphores

Mythili Vutukuru
CSE, IIT Bombay

Semaphore

- Another synchronization primitive like condition variables
 - Can be used to achieve similar synchronization between threads
- Semaphore is a variable with an underlying counter
 - Semaphore counter can be initialized to a suitable value
- Two functions on a semaphore variable
 - Down/wait decrements the counter by one, blocks the calling thread if the resulting value is negative
 - Up/post increments the counter by one, wakes up any one thread that is blocked on the semaphore
- Not possible to access counter value in any other way
 - For example, cannot check if counter is positive and only then call down

Example: use semaphore as a lock

- Consider semaphore variable “sem” initialized to value 1
- Multiple threads in a program can use semaphore for mutual exclusion
 - `down(sem)` //counter is 0, any further downs will wait here
 - critical section //accessed with mutual exclusion
 - `up(sem)` //counter incremented to 1, waiting thread woken up
- Such semaphore used as locks are called **binary semaphores**

T1

```
down(sem) //counter=0
critical section
....
...
up(sem)
```

T2

```
down(sem) //counter -1
(BLOCKED)

(woken by T1's call to up)
critical section
up(sem)
```

T1


```
down(sem) //counter 0
critical section
up(sem) //counter 1
```

T2

```
down(sem) //counter 0
critical section
up(sem) //counter 1
```

Example: T1 → T2

- Suppose we want two threads to synchronize as follows: T1 → T2
 - T1 does some work and only then T2 runs
 - If T2 starts before T1, it must wait until T1 finishes its task
- We can achieve this using a semaphore: “sem” initialized to 0
 - T1 does its work, calls up(sem)
 - T2 calls down(sem), then does its work



Semaphore implementation

- You may assume that up and down operations are implemented atomically (using locking internally as needed)
 - Counter accessed and updated with mutual exclusion
 - Need not worry about race conditions between up and down
 - No need to use extra locks to protect atomicity of wait/down
- But need to use separate locks or binary semaphores to access shared data in a program
 - Semaphore used for signaling doesn't provide mutual exclusion

Recap: Producer-consumer problem

- Producer and consumer threads, sharing data via a buffer of size N
 - Producers produce items, add into a shared buffer
 - Consumers consume item from shared buffer
- What kind of coordination is needed between threads?
 - Producer thread produces and places items into buffer, waits if the buffer is full → Consumer signals after making space in the buffer
 - Consumer thread consumes items from buffer, waits if the buffer is empty → Producer signals after producing items
- We have studied implementation with condition variables, similar implementation possible with semaphores

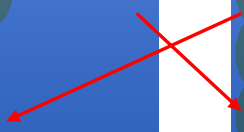


Producer-consumer using semaphores (1)

- One semaphore “sem_empty” is initialized to N, indicates number of empty slots in buffer available for producers to use
 - Producer does down every time it produces an item
 - Once all slots are filled, down operation blocks
 - Sleeping producer woken up by consumer that calls “up” after consuming
- Another semaphore “sem_filled” is initialized to 0, indicates number of filled slots in buffer that are ready to be consumed
- Easier solution than CV, no need to keep separate counter

```
//Producer  
down(sem_empty) //blocks if buffer full  
produce item  
up(sem_filled)//wakeup consumer
```

```
//Consumer  
down(sem_filled)//blocks if buffer empty  
consume item  
up(sem_empty)//wake up producer
```



Producer-consumer using semaphores (2)

- Note: semaphore solution does not use any locks by default
 - Condition variables had associated locks
- If buffer needs to be accessed correctly, needs extra locks via binary semaphores for mutual exclusion
- Use another semaphore mutex (initialized to 1) for locking

```
//Producer  
down(sem_empty) //blocks if buffer full  
down(mutex)  
produce item and add to buffer  
up(mutex)  
up(sem_filled)//wakeup consumer
```

```
//Consumer  
down(sem_filled)//blocks if buffer empty  
down(mutex)  
consume item and remove from buffer  
up(mutex)  
up(sem_empty)//wake up producer
```


Producer-consumer using semaphores: deadlock

- With condition variables, lock given to sleep/wait is released after the thread is safely put to sleep
 - No such concept of releasing any locks with semaphores
- With semaphore, if you do down with another binary semaphore/lock held, the lock will not be released on its own
- The solution shown below leads to deadlock: why?

```
//Producer
down(mutex)
down(sem_empty) //blocks if buffer full
produce item
up(sem_filled)//wakeup consumer
up(mutex)
```

```
//Consumer
down(mutex)
down(sem_filled)//blocks if buffer empty
consume item
up(sem_empty)//wake up producer
up(mutex)
```

Guidelines for using semaphores

- Semaphores can be used to do similar thread synchronization as CV
 - Waiting on CV ~ down operation on semaphore
 - Signaling on CV ~ up operation on semaphore
 - No equivalent of signal broadcast with semaphores
- Separate semaphores needed for signaling and mutual exclusion
- Semaphore **counter** can replace some integer/bool variables
 - But cannot access or change semaphore counter separately
- Careful with **deadlocks**
 - Ensure that “up” can run every time a thread blocks due to “down”
 - Note: no locks released when a thread blocks due to down operation
- Pay attention to **initial value** of semaphore

Example: Batched processing (1)

- Two kinds of threads in an application
 - Request threads, each containing an application request
 - Batch processor thread processes N requests at a time in a batch
- What kind of synchronization do we need?
 - Batch processing thread must wait until N requests arrive, then start batch
 - Request thread must wait until batch starts, then get processed and finish
- Example: suppose Covid-19 vaccination vial has 10 doses. Nurse waits for 10 patients to arrive, then opens the vial and vaccinates all 10

Example: Batched processing (2)

- Solution using two CVs: one for requests to wait, one for batch processor to wait
 - Other integer and boolean variables, mutex/lock for atomicity

```
//Request thread
```

```
lock(mutex)
```

```
count++
```

```
if(count == N)
```

```
    signal(cv_batch_processor)
```

```
while(not batch_started)
```

```
    wait(cv_request, mutex)
```

```
unlock(mutex)
```

```
//Batch processor thread
```

```
lock(mutex)
```

```
while(count < N)
```

```
    wait(cv_batch_processor, mutex)
```

```
    batch_started = true
```

```
    signal_broadcast(cv_request)
```

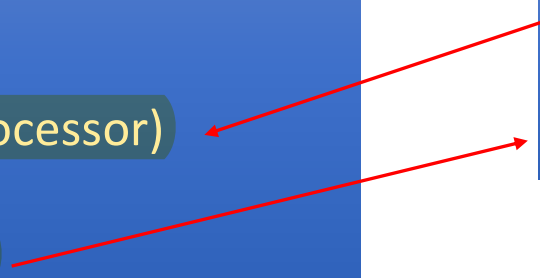
```
    unlock(mutex)
```

Example: Batched processing (3)

- Semaphore mutex initialized to 1, acts as lock to update count
- Semaphore sem_batch_processor, initialized to 0
 - Batch processor waits, until Nth request unblocks it
- Semaphore sem_request, initialized to 0
 - All N request threads wait on it (until batch starts)
 - When batch starts, batch processor thread does up N times to unblock all

```
//Request thread  
down(mutex)  
count++  
if(count == N)  
    up(sem_batch_processor)  
up(mutex)  
down(sem_request)
```

```
//Batch processor thread  
down(sem_batch_processor)  
//ready to start batch  
do N times: up(sem_request)
```



Example: Batched processing (4)

- Alternate pattern of solution
- Semaphore `sem_request`, initialized to 0
 - All N request threads wait on it (until batch starts)
 - When batch starts, batch processor does up once, unblocks only one thread
 - Each woken up request thread wakes up one other thread

```
//Request thread
down(mutex)
count++
if(count == N)
    up(sem_batch_processor)
up(mutex)
down(sem_request)
up(sem_request)
```

```
//Batch processor thread
down(sem_batch_processor)

up(sem_request)
```